



Utilization of Machine Learning (ML) in Prediction of House Pricing at the Real Estate Industry

Under the Supervision of:

Dr. Nik Hadiyan Binti Nik Azman

Prepared By:

Kermina Rafik Abd El Sayed

P-EM0248/23

Academic Year 2023/2024

School of Management

Universiti Sains Malaysia

Abstract

This study elucidates the transformative impact of the Fourth Industrial Revolution, resulting in the emergence of the property technology concept. This phenomenon has significantly affected and revolutionized traditional practices in the real estate industry, especially in the aftermath of the COVID-19 pandemic, with the integration of advanced technologies, including but not limited to artificial intelligence and machine learning. In the real estate industry, predicting house prices is an essential yet complex task. Achieving accurate predictions is made possible through machine learning technology, offering benefits to various stakeholders, including buyers, sellers, investors, and companies. The study explains and examines the significance of utilizing machine learning for accurate house price predictions, exemplifying the concept of property technology. This is demonstrated through an in-depth analysis of the Ames Housing dataset sourced from Kaggle. This involves detailed explanations of data exploration, data pre-processing, model building, and evaluation of models. The objective is to identify the model with the highest performance in predicting house prices. The study employs and compares multiple machine learning regression techniques using three main metrics which are R-squared, Mean Square Error and Root Mean Square Error, showcasing the practical implementation of house price prediction including Lasso, K-Nearest Neighbor Regressor, Support Vector Regressor, Decision Tree Regressor, Random Forest Regressor, Gradient Booster Regressor, eXtreme Gradient Boosting, LightGBM Regressor, and Voting Regressor. Our Champion Model was XGBoost regressor achieving highest R^2 squared, lowest MAE and lowest RMSE while the worst performing model was Decision Tree model with lowest R^2 squared, highest MAE and RMSE.

Key Words: Fourth Industrial Revolution, Property technology, Machine Learning, Real Estate, House Prices, Regression techniques, Lasso, KNN, SVR, DTR, RFR, GBR, XGBoost, LGBM, Voting Regressor.

Table of Contents

Abstract	I
Acknowledgment.....	V
Chapter 1	1
1.0 Introduction	1
1.1 Background of study	1
1.1.1 Property technology	1
1.1.2 Zillow Real Estate Website	2
1.1.3 House Price Prediction	2
1.1.4 main goal	3
1.1.5 Ames, Iowa City	4
1.2 Problem Statement.....	4
1.3 Study questions.....	5
1.4 Study objectives.....	5
1.5 Significance of the study.....	5
1.6 Scope of the study.....	6
1.7 Outline.....	6
Chapter 2	7
2.0 Literature Review	7
2.1 Fourth industrial revolution (Industry 4.0).....	7
2.2 Real estate, Industry 4.0, and Property technology	7
2.3 Utilizing machine learning in house price predictions.	8
2.3.1 Supervised machine learning.	8
2.3.2 Related work to house price prediction.	9
2.4 Regression machine learning models	11

2.5 Model Evaluation Techniques	13
Chapter 3	14
3.0 Methodology	14
3.1 Data Retrieval.....	15
3.2 Exploratory Data Analysis (EDA)	16
3.2.1 Basic Overview About the Data.....	16
3.3 Data Preprocessing	19
3.3.1 Handling Missing values	19
3.3.2 Log Transformation.....	20
3.3.3 Handling Outliers	21
3.3.4 Handling Categorical Features.....	24
3.3.5 Splitting the Data.....	24
3.3.6 Scaling and Splitting the Data.....	24
3.4 Modelling and Hyperparameter Tuning	24
Chapter 4	26
4.0 Findings.....	26
4.1 Hyperparameter tuning results	26
4.2 Model Performance.....	27
4.4 Discussion	30
Chapter 5	30
5.0 Conclusion.....	30
5.1 Limitations	31
5.2 Recommendations.....	31
References.....	33
Appendix	39

List Of Tables

Tabel 3. 1 Description of dataset features.....	15
Tabel 3. 2 Sale Price Descriptive Analysis	17
Tabel 3. 3 Top 10 Positive Correlated Features.....	18
Tabel 3. 4 Features Missing Values	19
Tabel 4. 1 Hyperparameter Tuning Values.....	26
Tabel 4. 2 Model Evaluation Comparison.	27

List of Figures

Figure 3. 1 Methodology workflow.....	14
Figure 3. 2 Sale Price Histogram.	18
Figure 3. 3 Sale Price Before Log Transformation.	20
Figure 3. 4 Sale Price After Log Transformation.	21
Figure 3. 5 Features with Outliers.....	22
Figure 3. 6 Features Without Outliers.	23
Figure 4. 1 Models R2 Scores.....	28
Figure 4. 2 Models MAE Scores.....	29
Figure 4. 3 Models RMSE Scores.....	29

Acknowledgment

First, I want to express my deepest gratitude for God presence, guidance, and support he offered me throughout my research work to complete successfully. Without God's help and support I would not be able to achieve anything. God led me through it all. I would like to thank Dr. Nik Hadiyan Binti Nik Azman , I was honored to work under her supervision. I am also extremely grateful for the unconditional love, support, encouragement and understanding that my parents, brother, and friends showed me throughout my master's journey. They played a huge part; without them I wouldn't be here. Finally, a special thanks to Eng. Basma for guiding and helping me whenever I needed throughout my project, she was always there to support me.

Chapter 1

1.0 Introduction

This chapter explains in detail the background of the study including the evolution of real estate industry, the importance of machine learning in house price prediction, followed by the problem statement discussed throughout this research paper, the research questions, the research objectives and finally the significance of the study, Scope of project and organization.

1.1 Background of study

Industry 4.0 (I4.0) also known as the fourth industrial revolution, which is initiated by digital technologies advancement, it transformed the economy and how industries operate through the integration of artificial intelligence (AI), machine learning, internet of things (IOT), large data, 3D printers, cloud computing and other advanced technology, leading to increasing automation, efficiency, and productivity, also encouraging data-driven decisions. Industry 4.0 is affecting how people live, work, and think. The usage of information technology and services became vital in the world of industries in order to be able to acquire a competitive advantage to differentiate a business from its competitors. (Machkour & Abriane, 2020)

1.1.1 Property technology

The real estate industry could not avoid the reshaping force of industry 4.0. (Starr et al., 2020) Especially after the covid-19 pandemic that led to changing the rules of real estate due to the need for remote and limited interactions where real estate companies and industry needed to quickly adapt to modern technologies. Property technology or prop-tech concept has emerged as a subset of industry 4.0 which simply means the use of digital technology in the real estate industry. Prop-tech is disrupting the practices of traditional real estate market in terms of the way properties are bought, sold, and financed through the presence of technological innovation applications that are transforming the real estate industry which include new category of products, websites, and mobile apps designed to streamline the process of browsing, buying, and building real estate. (cycles & Text, n.d.)

Prop-tech includes but not limited to artificial intelligence and machine learning that are used in predicting house prices, blockchain technology that reduces fraud and enhances the security and

transparency of real estate transactions, and data analytics that aids professionals to gain insights from large real estate datasets enabling for better investments decisions and risk assessments.

These technologies are transforming and enhancing various aspects of the real estate industry. (Putatunda, 2019) (Saiz, 2020). Prop-tech investments are continuously growing where most of the investment is concentrated in the United States, China, and Asia specific (APAC) where USA investments reached \$ 61.1 billion in the first half of 2022. while China and India reached 12.5 billion and 9.1 billion, respectively. In Europe, the top investors are in Spain, the UK and Germany. Prop-tech investments are expected to grow and consider as an opportunity for real estate companies (cycles & Text, n.d.).

1.1.2 Zillow Real Estate Website

Zillow is well-known in the real estate online market. It is one of the most visited real estate websites in the United States. Zillow company is a successful example of a prop-tech company that utilizes modern technology in their processes. Zillow offers various on-demand services for their customers such as property price listing, home valuation tools and market data to help them in selling, buying, renting, and financing properties with transparency. One of the most important tools on their platform is “Zestimate” which is an automated home valuation model that forecast home’s market value considering home characteristics, amenities, location, and market trends such as supply and demand, interest rates and economic conditions. Zillow company is challenging the traditional real estate market by providing more options and information to the buyers and sellers giving them the opportunity to make better informed decisions. (Zillow, n.d.)

1.1.3 House Price Prediction

In this study the main focus will be on utilization of machine learning in prediction of house prices in the real estate industry as an application of property technology with highlighting that property technology and machine learning as an opportunity for companies and real estate stake holders. Housing is fundamental for everyone whether for investing or for living in, and house price prediction is a very complex task. house prices are affected by countless number of features including external features such as the physical characteristics of the building, its location, and the neighborhood and the internal features such as the square feet , number of rooms, condition of kitchen, and living area (Kumar & Syed, 2023). Experts in real estate such as agents rely on years of experience for the valuation of houses and properties (Geerts et al., 2023). On the other hand,

investors and buyers have multiple concerns when buying a house such as the correct cost for the houses and whether the features of these houses properly synchronize with the price. Thus, it becomes essential to accurately determine the prices of the house or property.

House prices can be predicted using machine learning models through analyzing datasets of historical sales data that contain house features, demographic information, and other indicators on which the price depends in order to uncover hidden patterns and relationships that may not be obvious to traditional methods. By performing predictive analysis of future trends potential scenarios can be predicted and used to take important strategic decisions. Building a model for houses prediction can benefit both the house sellers and buyers in making well-informed decisions, lowering costs, and increase profit. (Kansal et al., 2023)

In house price prediction, regression techniques are often used among other options where a single target is used as the output which in our case is the price and one or more predictors are used as input indicators which are the house features (Yadav et al., 2023). There are various regression models that can be used including linear regression, random forest regressor, support vector machine regressor, k-nearest neighbor and others as everyday machine learning is evolving to find better and more efficient solutions.

1.1.4 main goal

This study aims to build several machine learning regression techniques to explain the process of house price prediction through machine learning by utilizing a comprehensive Ames housing dataset from Kaggle website. The study will compare these models' performance to select the champion model that can achieve higher performance in Ames, Iowa house price predictions with considering that there is no one size fits all model, best machine learning model to use depends on the size and complexity of the data.

1.1.5 Ames, Iowa City

Before diving into this study, it is important to know more about Ames, Iowa city. Ames is in the state of Iowa, which is a growing city located in United States of America with a population of more than 65,000, this number includes students attending Iowa State University who reside in Ames during the school year or longer. As a growing city, Ames continues to focus on building a strong community filled with opportunities for all citizens. Also, the city is expanding housing opportunities and diversifying housing choices. The city has its own assessor's office, where the assessor's primary duty is to assess all real property. The real property is revalued every two years. As an on-going process of gathering and reviewing information, measuring, and listing new construction and investigating sales of real estate to determine the fair market value of property (*City of Ames, IA | Home*, n.d.).

1.2 Problem Statement

It is a critical issue for Ames, Iowa real estate stakeholders including agencies, and local government officials in the real estate industry to accurately predict fair house prices as it is a growing population city and evaluation of houses is essential. However, house prices prediction is not an easy task especially when using traditional methods of comparing different houses with various different features that affect prices. These traditional methods are time consuming and often inaccurate as house prices are influenced by several factors such as location, size, condition, , amenities and many others. False price prediction may lead to several problems. If real estate agencies price houses too high this could challenge the selling process and cost them huge monetary loss, while pricing houses too low means not getting the full value, also buyers might end up paying too much or lose out the opportunity of good deals. On the other hand, investors could make wrong investing decisions. Accordingly, false prices can disrupt the real estate market and damage public trust. Furthermore, Banks can face credit risk such as risk related to mortgage loans also the government represented in tax authority may face losses during collecting real estate taxes. So, ensuring accurate house pricing is essential to avoid all the previous issues in the real estate market. Addressing previous issues, machine learning (ML) can be used to solve this problem. Specifically tailored to the Ames housing market, a machine learning model can be trained on Ames housing dataset to learn the relationship between the factors that affect house prices. Once trained, these machine learning models can be used to predict the prices of new houses

with a high degree of accuracy, saving effort, cost, time, and can help real estate stakeholders in Ames, Iowa city to predict house prices and to avoid previous mentioned issues leading to saving resources. As well as making more informed decisions regarding house prices.

1.3 Study questions

- I. Which regression model achieves the best performance in Ames house price prediction?
- II. What are the major features that significantly influence prediction of house prices?
- III. What are the most effective data preprocessing techniques for Ames house price prediction?
- IV. Which of the feature engineering techniques can be utilized to enhance models' prediction powers?
- V. What are possible future work and limitations?

1.4 Study objectives

- I. To evaluate the performance of different regression models in predicting house prices in Ames housing dataset from Kaggle.
- II. To understand different features that influence house prices prediction.
- III. To identify effective preprocessing techniques for Ames house prices prediction.
- IV. To identify the most suitable feature engineering technique to enhance models' prediction.
- V. To highlight possible future work and limitations.

1.5 Significance of the study

Nowadays, with the presence of modern technology, accurately predicting house prices has become crucial in the real estate market for many stakeholders, specifically in Ames, Iowa. Successfully implementing house price prediction models has several benefits. Accurate predictions that can help individuals to make more informed decisions concerning property investments, purchases, and sales. Also, reliable price estimates can encourage transparency in the real estate market and reduce uncertainty. It can help in assessing and managing real estate investment risks. Finally, the insights gained from this project can encourage the development of new products and services tailored to Ames, Iowa real estate market.

1.6 Scope of the study

The study focuses on building and comparing multiple regression machine learning techniques for predicting house prices in Ames , Iowa city. In addition to comparing the performance of these models using multiple evaluation techniques. This study does not include any further deployment of machine learning.

1.7 Outline

This project paper consists of five main chapters which are **Chapter 1** the introduction: this chapter is an overview of the project topic which is “ Utilization of machine learning (ML) in prediction of house price at the real estate industry.”. It starts with an overview of fourth industrial revolution and its impact on the real estate industry through the emergence of Property technology concept and the importance of house price prediction using machine learning, highlighting the problem statements, the project objectives, project questions, and significance of the study. **Chapter 2** the literature review, this chapter reviews and discusses the previous work regarding the topic of house price predictions using machine learning. **Chapter 3** the methodology, which will include explaining the data collection and source, the data preprocessing, the data exploratory and the machine learning model used. **Chapter 4** the findings and discussion from the data analysis and applied machine learning models. Finally, **Chapter 5** will discuss the limitations, recommendations, and conclusions reached.

Chapter 2

2.0 Literature Review

This literature review builds a strong foundation to support this study which is achieved through reviewing past work regarding real estate industry including the explanation of fourth industrial revolution, property technology and its application, in addition to emphasizing the role of machine learning in house price prediction which can benefits stakeholders such as house buyers , seller and investors while describing various used machine learning supervised models and techniques.

2.1 Fourth industrial revolution (Industry 4.0)

Fourth industrial revolution, what is known as industry 4.0 was first defined by Klaus Schwab, founder of the world economic forum which described the industry 4.0, showing a society where people transition through technology between digital and physical world to facilitate their daily activities (Xu et al., 2018). Moreover, as the fourth industrial revolution supports digital technology it is also described as “The Digital Revolution”. This technology includes the combination of artificial intelligence, robotics, 3D printing, internet of systems (IOS), internet of things (IOT), cloud computing and large data. The main goal of industry 4.0 is to improve peoples’ standard of living and to increase profits through being more faster, effective, and efficient leading to better industrial transformation (Dogaru, 2020) (Machkour & Abriane, 2020). The adoption of industry 4.0 technologies in companies and industry became important as it provides them with innovation, competitiveness, and growth powers (Bai et al., 2020).

2.2 Real estate, Industry 4.0, and Property technology

Real estate as a valuable asset and industry, couldn’t avoid the huge impact made by the fourth industrial revolution (industry 4.0) innovations, and prop-tech has emerged, short for property technology which means the digital transformation in real estate industry (Baum et al., 2020). Prop-tech is a movement changing the shape and business model of the real estate industry including innovative technology and new products such as virtual reality (VR), Artificial intelligence (AI) and machine learning, building information modeling (BIM), internet of things (IoT), Blockchain, in addition to smart cities and homes which revolutionized the real estate industry. Property technology led to the development of smart real state which is the technology platforms for operating and managing real estate processes (GÓRSKA et al., 2022).

Prop-tech proved to be a key point for real estate companies to acquire competitive advantage in the industry through the use of these innovative technologies. Key players real estate companies such as Zillow, Airbnb and WeWork proved that the impact of such business models can be huge (Bittini et al., 2022), these companies have gained the market share from traditional real estate agents by providing straightforward information aggregation proposition in addition to facilitating contractual offers for properties on their platforms. According to a research study the global prop-tech market share in terms of revenue was valued at \$19.5 billion in 2022 and it is expected to reach around \$32.2 billion by 2030 with a growth rate of 6.5% between 2023 to 2030 (Research, n.d.)(Yahoo Finance, 2023). The valuation process of properties and houses can be time consuming leading to several drawbacks such as wasting time, increase of uncertainty and delays in the sale processes. Prop-tech companies solved these problems through the utilization of machine learning for prediction individual property price which helps in reducing property selling and purchasing needed time in addition to increasing market transparency, and decreasing transaction costs which benefits both buyers, sellers and investors (Baum et al., 2020).

2.3 Utilizing machine learning in house price predictions.

Machine learning as subfield of artificial intelligence and an example of property technology proved to be a successful method for predicting house prices accurately as the traditional methods are complex for estimating house prices due to the presence of various of factors that can affect the valuation of house prices such as amenities, number of rooms, house conditions, number of bathrooms, condition of kitchen, and garages. Machine learning algorithm approved to be able to find hidden patterns and relationship addressing the most important house factors that affect house prices leading to accurate prediction of house prices. Accurate prediction of house prices has various benefits including helping buyers and sellers to make an informed decision regarding house prices (Kansal et al., 2023).

2.3.1 Supervised machine learning.

In supervised machine learning the training data set contains a target or desired solution which needs to be predicted according to understanding different features patterns. The supervised machine learning is classified into classification type such as prediction of emails whether it belongs to spam class or ham class, the other supervised type is called regression, and it is used

when the target which needed to be predicted is numeric value. In this project regression techniques are used where the target is house prices and predictors are different house features (Géron, 2022).

2.3.2 Related work to house price prediction.

In (Truong et al., 2020) research, the authors applied both traditional and advanced machine learning techniques to examine the difference between several advanced models. Using dataset called “Housing Price in Beijing” containing 300,000 observations and 26 variables, they built and evaluated the following models including Random Forest, XGBoost, LightGBM, hybrid regression and stacked Generalization Regression. After data preprocessing, data splitting, model building and evaluation of models using Root Mean Square Logarithmic Error (RMSLE), the result showed that all of the method showed a desired outcome but each method has its own pros and cons, random forest had high time complexity and was prone to overfitting, XGBoost and LightGBM performed well in terms of time complexity, finally hybrid regression and stacked generalization showed high accuracy however they had high time complexity. The study suggests considering fast ways for complex models.

In (Abdul-Rahman et al., 2021) research they utilized advanced machine learning models for house prices prediction comparing two recent models, LightGBM and XGBoost using combined data set called “Property Listing in Kuala Lumpur” from Kaggle and Google Map with 21984 observation and 11 feature, the study used root mean square error (RMSE), mean absolute error (MAE), and adjusted r^2 value. They found that XGBoost outperformed LightGBM. They proven the reliability of the model by deploying it using another sample of data and the results showed a small variance between the actual and predicted price.

In (Ho et al., 2021) paper, authors used three machine learning techniques for house price prediction including support vector machine (SVM), random forest (RF), and gradient boosting machine (GBM). The models were applied on Hong Kong dataset. The result showed that (RMSE), absolute percentage error (MAPE), of Rf and GBM have achieved better performance than SVM model. However, the study admitted that SVM is useful in data fitting as it produces accurate prediction in tight time constraints. In conclusion the study agreed on the promising opportunity of machine learning property valuation, it also highlighted the importance of feature selection.

In (Hanuma Reddy & Sriramy, 2022) this study used and predicted real estate prices comparing performance of two model which are voting regressor and linear regression models. For the voting regression model the author combined three regression models to obtain final result including linear regression, decision tree, and random forest regression. After evaluating and training the models, the results showed that the voting regression outperformed the linear regression and achieved higher accuracy.

The authors (Zhou et al., 2023), focused on building effective models by comparing several models such as linear regression, random forest regressor, XGBoost, support vector machine (SVM) regressor, K-Nearest Neighbor (KNN), and linear regression. The results showed that Random Forest Regressor outperformed other models with R2 score 99.97% while the least performing model was SVM with R2 score -4.11%, the champion model for house prices prediction was Random Forest.

In (Kumar & Syed, 2023) research, they highlighted the importance of price prediction using machine learning and how a predicting model can benefit buyers to find fair purchase price as well as sellers to find the suitable selling house prices. They also used dataset of six different cities from Kaggle, consisting of 40 features for house price prediction. They developed and compared six machine learning models based on their accuracy including linear regression, random forest , decision tree, KNN , XGBoost, and SVM. The result showed that XGBoost outperformed the other five models with the highest accuracy.

In (Yadav et al., 2023), the paper aim was to examine the main characteristics of real estate property affect price and provide insights for each characteristics. In this study several machine learning models were built such as linear regression, Lasso regression, Decision tree, and GridSearchCV to predict costs of houses using Bangalore dataset from Kaggle. The study focused on Linear regression results, which showed a score of 85% which was okay, then K-fold cross validation was used for optimality giving score 80%. The study results showed that linear regression had the highest prediction accuracy than other methods and recommended applying other methods in the future such as ensemble techniques for better accuracy.

In (El Mouna et al., 2023) study, they addressed house price predictions by applying three machine learning method including linear regression (LN), Random Forest (RF) and GradientBoosting (GB) on Melbourne real estate dataset which includes 34,857 sales and 21 house features. The

models were evaluated using (RMSE), (MAE) and R- square. The results showed that Gradient Boosting outperformed the other two models. It calculated maximum R-square explaining most of variability in the dataset with 0.828. in addition, it calculated lower error values for RMSE and MAE.

In (Sibindi et al., 2023) study, as Boosting ensemble techniques are widely used in house price predictions, they focused on developing hybrid LGBM and XGBoost model to prevent overfitting through minimizing variance and improving accuracy the hybrid model was compared to LGBM, XGBoost, Adaboost, and GBM models and the models were evaluated using Mean Square Error (MSE), Mean Average Error (MAE), and Mean Absolute Percentage Error (MAPE). The results showed that the hybrid model outperformed other models achieving 0.193, 0.285 and 0.156 respectively.

The above literature review gives an overview of different machine learning algorithms and techniques utilized in prediction of house prices in the real estate industry on various datasets with different sizes and features. In addition to explaining the evaluation methods used for evaluating the performance of these models.

2.4 Regression machine learning models

There are various types of machine learning models, but this study focuses on the following machine learning regression models which are widely used in house price predictions.

Least absolute shrinkage and selection operator regression (LASSO):

LASSO is regularized version of linear regression. A unique characteristic of LASSO is that it eliminates the weight of less important features. It automatically performs features selection (Géron, 2022).

K-nearest neighbor regressor (KNN):

KNN was first invented by Evelyn Fix and Joseph Hodges in 1951. It is a supervised machine learning technique and one of the easiest algorithms. KNN uses similarity measures, it simply assumes that similar data points are near to each other. K in KNN refers to the number of neighbors of the new data points. It is essential to choose the accurate value of K as low value of K can lead to noisy results while high value can create confusion, although there is no fixed value for K, the

standard is 5. In KNN The distance between data points is often measured using Euclidean distance. From KNN pros, is that it is easy and fast to apply and it is resistant to noise in the training data, however it is tricky to decide the accurate value of K and due to the calculation for the distance between all data points, it leads for high cost of computation (Bansal et al., 2022).

Support Vector Regressor (SVR):

SVR is a powerful model which has the ability to perform linear and nonlinear regressions. It works best with small to medium sized nonlinear datasets (hundreds to thousands of observations). However, it doesn't fit well with very large datasets. The main aim of SVR is to find the most suitable decision boundaries which is known as hyperplane that separates the data points while limiting margin violations (Géron, 2022) (Bansal et al., 2022).

Decision Tree Regressor (DTR):

Decision Tree is a powerful algorithm that is capable of fitting complex datasets. It is also a fundamental component to other machine learning algorithms such as random forest and XGBoost models. It has a tree-like structure where internal nodes represent a decision based on a feature and each leaf is a predicted outcome (Kansal et al., 2023) (Géron, 2022) (Vemuri, 2020).

Gradient Boosting Regressor (GBR):

It is a well-known boosting algorithm which is one of ensemble approaches that involves combining and training several weak learners (simple models) and capitalizing them into a strong learner (stronger model). Gradient Boosting works in a sequential manner in which each added predictor to the ensemble correct its predecessor through trying to fit the new predictor to the residual errors made by the previous predictor (Géron, 2022) (Vemuri, 2020).

Random Forest Regressor (RFR):

RFR is another ensemble algorithm that trains and combines multiple decision trees which resemble a forest. Each decision tree is trained on a unique dataset. Furthermore at each decision tree node, different subset of features are used for splitting which result in minimizing the influence of any single feature leading to the overall robustness of the model and achieving a high performance (Alyami et al., 2024) (Breiman, 2001) (Géron, 2022).

Extreme gradient boosting (XGBoost):

XGBoost is a tree boosting system. It is an ensemble learning technique that combines several decision trees. It is also known for its speed and performance when working with large datasets that consist of complex features. The model is designed to improve accuracy and strength of the model (Chen & Guestrin, 2016).

Light Gradient Boosting (LightGBM):

LightGBM is a gradient boosting decision tree model designed for large datasets and high dimensional features. It is designed to have the following advantages including faster training and prediction speed, higher efficiency, low memory usage, and higher accuracy (Ke et al., 2017).

Voting Regressor (VR):

Voting Regressor is another ensemble technique which combine and train numerous regression models for the purposes of aggregating the predictions of each one of the base models and produces a final prediction based on vote majority that results in an improved performance (Hanuma Reddy & Sriramya, 2022).

2.5 Model Evaluation Techniques

There are multiple performance measures for evaluating machine learning models that can be used to evaluate models, which are:

Root Mean Square Error (RMSE)

It is the most preferred performance measure for regression problems. It gives a good idea of how much error is represented in the model predictions by giving higher weights to large errors. However, if there are many outliers it's preferable to use another measure technique (Géron, 2022).

Mean Absolute Error (MAE)

It is a more suitable performance measure when there are many outliers as it is more sensitive than RMSE. It measures the difference between actual and predicted values (Géron, 2022).

R-Squared

It returns the coefficient of determination; it explains how much variability in the target variable is explained by the features variables. R-squared range between 0 to 1, the closer R^2 to 1 the more of the variance in the target variable is explained by the features (Albon, 2018).

Chapter 3

3.0 Methodology

The study's methodology and steps are explained in this chapter. The steps were performed using 'Google Collab' which is a tool that offers free access to Jupyter Notebook environment on cloud. In this study the whole process was conducted using python programming language. The methodology process starts from data retrieval, data exploration ,data pre-processing, feature engineering and finally, building and training models to figure out the champion model for predicting house prices including Least Absolute Shrinkage and Selection Operator (Lasso), K-Nearest Neighbor Regressor (KNN) , Support Vector Regressor (SVR) , Decision Tree Regressor (DTR) , Random Forest Regressor (RFR), Gradient Boosting Regressor (GBR), eXtreme Gradient Boosting (XGBoost) , Light Gradient Boosting Machine Regressor (LGBM), Voting Regressor.

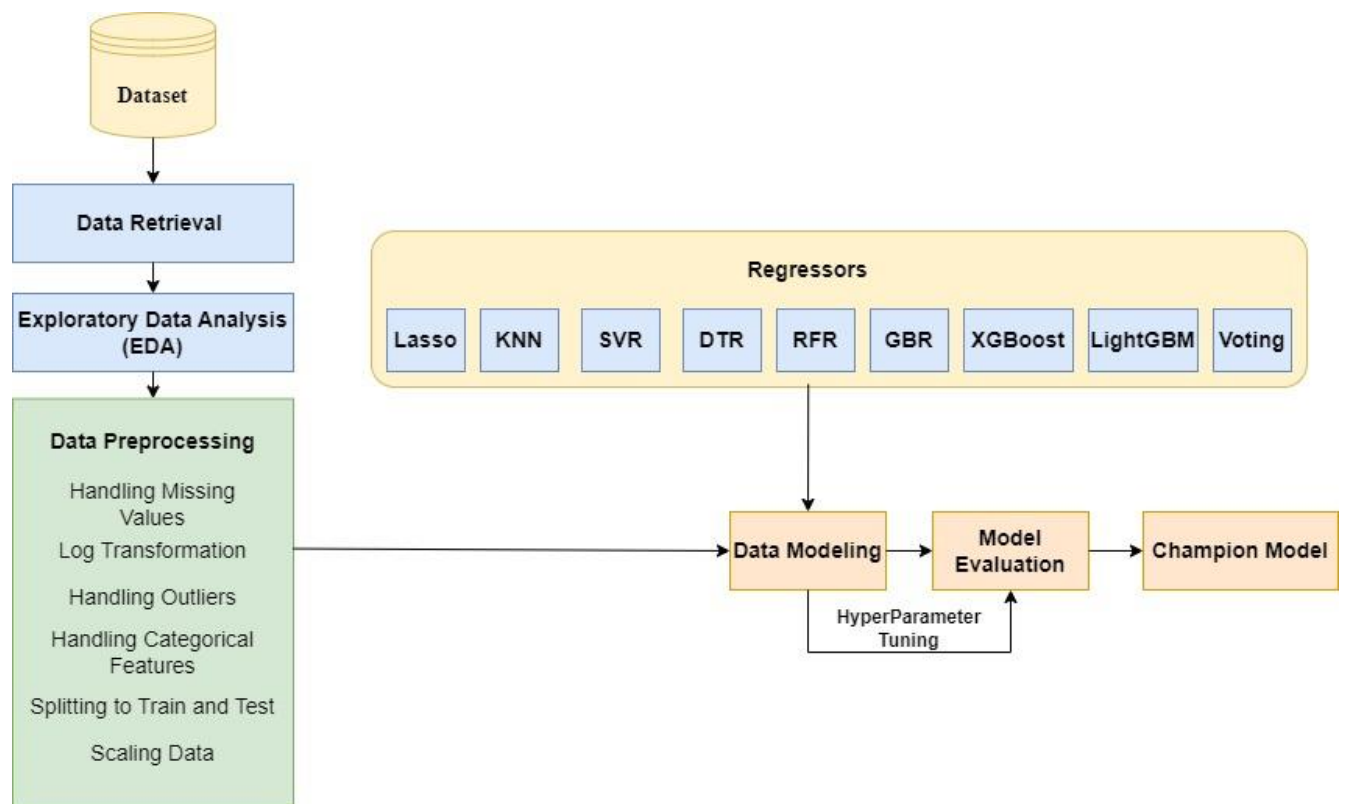


Figure 3. 1 Methodology workflow.

3.1 Data Retrieval

The dataset used for this study is called ‘Ames Housing dataset’ that was downloaded from Kaggle website in the form of comma-separated value (CSV) file format. The data set contains home sales that occurred in Ames, Iowa between 2006 and 2010 with 1460 observations and 79 explanatory features (categorical and numerical) describing residential houses in Ames, Iowa along with each house’s sale price as the target variable. The data was imported to python using **read_csv()** which is used to read csv files. The following table contains descriptions for each feature included in the dataset.

Tabel 3. 1 Description of dataset features.

Feature	Description	Feature	Description
• MSSubClass	Identifies the type of dwelling involved in the sale.	• MSZoning	Identifies the general zoning classification of the sale.
• LotFrontage	Linear feet of street connected to property.	• LotArea	Lot size in square feet.
• Street	Type of road access to property.	• Alley	Type of alley access to property.
• LotShape	General shape of property.	• LandContour	Flatness of the property.
• Utilities	Type of utilities available.	• LotConfig	Lot configuration.
• LandSlope	Slope of property.	• Neighborhood	Physical locations within Ames’s city limits.
• Condition1	Proximity to various conditions.	• Condition2	Proximity to various conditions.
• BldgType	Type of dwelling.	• HouseStyle	Style of dwelling.
• OverallQual	Rates the overall material and finish of the house.	• OverallCond	Rates the overall condition of the house.
• YearBuilt	Original construction date.	• YearRemodAdd	Remodel date.
• RoofMatl	Roof material.	• Exterior1st	Exterior covering on house.
• Exterior2nd	Exterior covering on house.	• MasVnrType	Masonry veneer type.
• MasVnrArea	Masonry veneer area in square feet.	• ExterQual	Evaluates the quality of the material on the exterior.
• ExterCond	Evaluates the present condition of the material on the exterior.	• Foundation	Type of foundation.
• BsmtQual	Evaluates the height of the basement.	• BsmtCond	Evaluates the general condition of the basement.
• BsmtExposure	Refers to walkout or garden level walls.	• BsmtFinType1	Rating of basement finished area.
• BsmtFinSF1	Type 1 finished square feet.	• BsmtFinType2	Rating of basement finished area.
• BsmtFinSF2	Type 2 finished square feet.	• BsmtUnfSF	Unfinished square feet of basement area.

• TotalBsmtSF	Total square feet of basement area.	• Heating	Type of heating.
• HeatingQC	Heating quality and condition.	• CentralAir	Central air conditioning.
• Electrical	Electrical system.	• 1stFlrSF	First Floor square feet.
• 2ndFlrSF	Second floor square feet.	• LowQualFinSF	Low quality finished square feet (all floors).
• GrLivArea	Above grade (ground) living area square feet.	• BsmtFullBath	Basement full bathrooms.
• BsmtHalfBath	Basement half bathrooms.	• FullBath	Full bathrooms above grade.
• HalfBath	Half baths above grade.	• Bedroom	Bedrooms above grade.
• Kitchen	Kitchens above grade.	• KitchenQual	Kitchen quality.
• TotRmsAbvGrd	Total rooms above grade.	• Functional	Home functionality.
• Fireplaces	Number of fireplaces.	• FireplaceQu	Fireplace quality.
• GarageType	Garage location.	• GarageYrBlt	Year garage was built.
• GarageFinish	Interior finish of the garage.	• GarageCars	Size of garage in car capacity.
• GarageArea	Size of garage in square feet.	• GarageQual	Garage quality.
• GarageCond	Garage condition.	• PavedDrive	Paved driveway.
• WoodDeckSF	Wood deck area in square feet.	• OpenPorchSF	Open porch area in square feet.
• EnclosedPorch	Enclosed porch area in square feet.	• 3SsnPorch	Three season porch area in square feet.
• ScreenPorch	Screen porch area in square feet.	• PoolArea	Pool area in square feet.
• PoolQC	Pool quality.	• Fence	Fence quality.
• MiscFeature	Miscellaneous feature not covered in other categories.	• MiscVal	Value of miscellaneous feature.
• MoSold	Month Sold (MM).	• YrSold	Year Sold (YYYY).
• SaleType	Type of sale.	• SaleCondition	Condition of sale.

3.2 Exploratory Data Analysis (EDA)

Exploratory Data Analysis is an essential step for every project and study. It was applied for exploring, investigating, and summarizing data to uncover patterns and trends to understand data fundamentals. It is the step preceding advanced analysis and modelling which gives a basic view about the data. (*Pandas Documentation — Pandas 2.1.4 Documentation*, n.d.)

3.2.1 Basic Overview About the Data

In order to have a quick understanding of the dataset we implemented the following steps:

We used **house_df.head()** function to retrieve the first five rows of the data, **house_df.columns** to give the complete lists of column names, **house_df.shape** function to retrieve the number of rows and columns of the data, we used all the previous to have a quick overview of the data

columns and observations. Then, we used **house_df.duplicated().sum()** function to check for possible duplicates represented in the data.

The Id column is useless for our study, we dropped it using the following function **house_df.drop(['Id'], axis = 1, inplace = True)**.

To further explore the data, we used **house_df.info()** function to return comprehensive information about the data such as names, non-null count, and data type of each column. Additionally, we used **house_df.describe().T** that generates descriptive analysis for the data including count of observations, mean, standard deviation, minimum, maximum, 25% quartile, 50% quartile, and 75% quartile for each feature in the dataset to understand central of tendency, dispersion, and shape of data. In addition to using the **describe()** function for the whole dataset, it was used to focus on our target variable 'SalePrice' as well **house_df['SalePrice'].describe()**. The mean (180921.195890) resulted from SalePrice descriptive analysis is higher than the median (163000.000000) which suggests a right skewed distribution, so we needed to visualize it later in this chapter. The following schedule explains all descriptive analysis values for SalePrice.

Tabel 3. 2 Sale Price Descriptive Analysis

Measure	values
Count	1460.000000
Mean	180921.195890
STD	79442.502883
Min	34900.000000
25%	129975.000000
50%	163000.000000
75%	214000.000000
max	755000.000000

Furthermore, to understand our target distribution, we visualized it using histogram plot function **sns.histplot(house_df['SalePrice'])** to reveal its skewness pattern that was right skewed and will be altered later in this chapter, it is represented by figure 3.2. We also measured its variance using **house_df['SalePrice'].var()** function to understand how much variability exists in the target variable.




Figure 3. 2 Sale Price Histogram.

Then, we retrieved the numerical and categorical features names, they were 37 numerical features and 43 categorical variables. We used horizontal bar plot **sns.barplot()** to visualize all categorical features against SalePrice to clearly understand their values and their unique classes. We also visualized numerical features using pair plot function **sns.pairplot()** that showed features outliers that ought to be removed in addition to their relationship with each other and with the SalePrice. We used heatmap **sns.heatmap()** function to understand numerical features and their relationship and correlations with SalePrice feature, as well as calculating correlation matrix using **corr()** function, correlation ranges between 1 and -1 the closest the correlation to 1 the strongest. The following table contains the top 10 positive correlated variables with SalePrice feature including OverallQual, GrLivArea, GarageCars, GarageArea, TotalBsmtSF, 1stFlrSF, FullBath, TotRmsAbvGrd, YearBuilt, and YearRemodAdd which are highly correlated with SalePrice. We can find that overall house quality has the strongest correlation with house sale price.

Table 3. 3 Top 10 Positive Correlated Features

Feature	Positive Correlation
OverallQual	0.790982
GrLivArea	0.708624
GarageCars	0.640409
GarageArea	0.623431

TotalBsmtSF	0.613581
1stFlrSF	0.605852
FullBath	0.560664
TotRmsAbvGrd	0.533723
YearBuilt	0.522897
YearRemodAdd	0.507101

3.3 Data Preprocessing

After retrieving the data and performing exploratory data analysis to get initial insights, a data preprocessing was performed to prepare the data to be in a suitable format for the machine learning models.

3.3.1 Handling Missing values

Most machine learning models cannot work with features that contain missing values so, they need to be handled, in order for machine learning models to work accurately. We calculated the null values for every feature, with their percentages and their data type in order to be able to see the full picture. We found that 19 features out of 79 contained null values, and they were handled differently according to the number of null values, their data type, and the data description, we didn't want to remove any information as the data is relatively small and according to real estate every information should be considered. The following table represents features that had null values and how they were handled.

Tabel 3. 4 Features Missing Values

Feature	No. Nulls	Replaced with	Feature	No. Nulls	Replaced with
PoolQC	1453	None (no pool)	GarageType	81	None (no garage)
MiscFeature	1406	None (no extra feature)	GarageFinish	81	None
Alley	1369	None (no alley)	GarageQual	81	None
Fence	1179	None (no fence)	BsmtExposure	38	None (No basement)
FireplaceQu	690	None	BsmtFinType2	38	None

		(no fireplace)			
LotFrontage	259	Median	BsmtCond	37	None
GarageYrBlt	81	Median	BsmtQual	37	None
GarageCond	81	None	BsmtFinType1	37	None
MasVnrArea	8	Median	MasVnrType	8	Mode
Electrical	1	Mode			

3.3.2 Log Transformation

In our EDA the target variable (SalePrice) was heavily skewed to the right and many machine learning models assume normal distribution of the target variable. Accordingly, the target variable was log transformed using **np.log()** NumPy function, so that it can be normally distributed or close to normal. We can use **np.exp()** later after prediction to reverse the transformation for predicted values. The following figures represent target variable before and after transformation.




Figure 3. 3 Sale Price Before Log Transformation.




Figure 3. 4 Sale Price After Log Transformation.

3.3.3 Handling Outliers

Outliers are data points which are extremely high or low and deviate from other data points that result from data entry error, computation error or natural outliers and can lead to skewed and inaccurate predictions. We defined a range of acceptable values between the 1st and 3rd quartile by using interquartile range (IQR) value. The values that fall outside the range are considered outliers and they can be treated either by removing or using a capping method. In this study we used capping methods. The capping method was used to handle outliers through setting threshold value and any outlier that is beyond the threshold is replaced with a specific value, as a result it preserved information of the outlier instead of removing them. We visualized each numerical feature using Boxplot before and after using capping method, for checking the capping effect on features that contained outliers. The following Figures represent outliers before and after handling them.

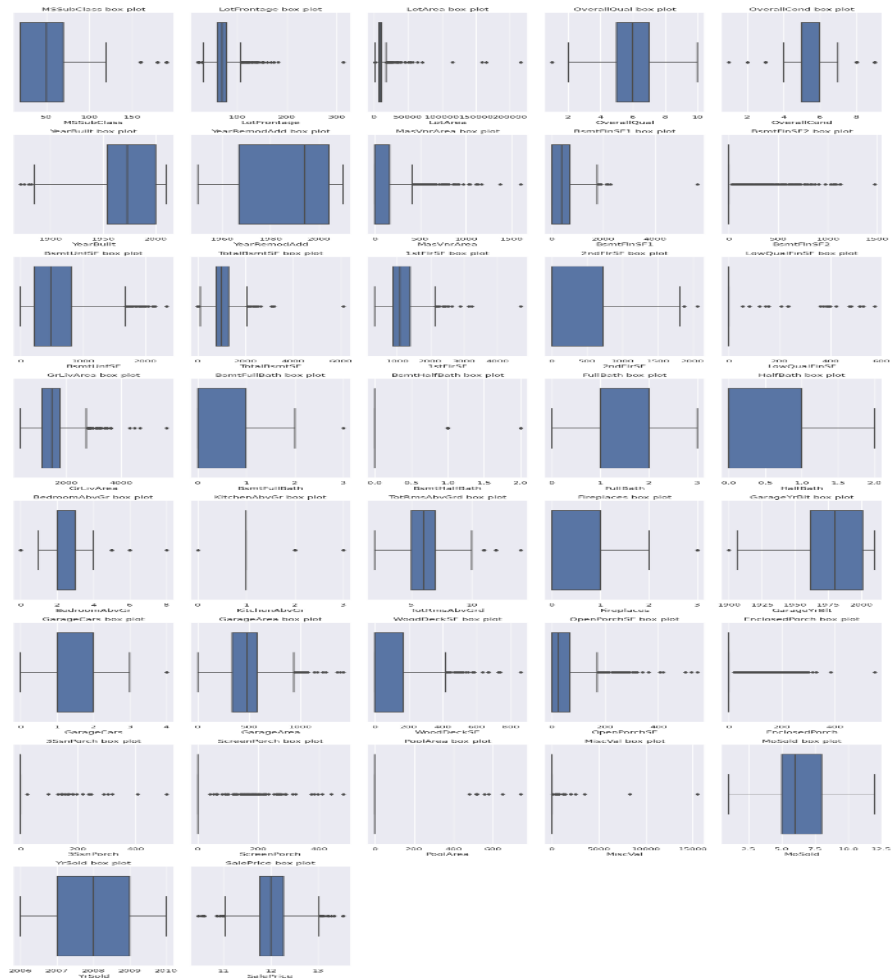


Figure 3. 5 Features with Outliers.



3.3.4 Handling Categorical Features

Since, our dataset contains a lot of categorical features and most machine learning models can only take numerical values as input, they must be encoded into numerical values before being used to test and train the models. We used **LabelEncoder()** from `sklearn.preprocessing` library to encode categorical features where each category within the feature is assigned to a unique integer starting from 0 to number of classes.

3.3.5 Splitting the Data

Before building models, the dataset needs to be split into train set and test set. The train set is used to train and fit the model while test set is used to evaluate the model performance. In this study the dataset was split into 80% for model training and 20% for model testing. We split the data using **train_test_split()** to split the data into `X_train`, `X_test`, `y_train`, and `y_test` where `X` was the features with dropping the target and `y` was the target.

3.3.6 Scaling and Splitting the Data

Feature Scaling is one of the most important transformations that need to be applied to our data. When the data have very different scales, machine learning models cannot perform well so they need to be transformed. In this study we used **StandardScaler()** for scaling our features. It is important to fit `StandardScaler` to the training data only and not the testing data in order to prevent data leakage which is the transferring of information from training set to testing set.

3.4 Modelling and Hyperparameter Tuning

In this study, we built nine different machine learning models in order to compare them and choose the highest performance model to be used for house price predictions including Lasso, KNN, SVR, DTR, RFR, GBR, XGBoost, LGBM, Voting Regressor. We used **cross_val_score()** function on each model in order to evaluate models' performance through cross validation which divide training data into multiple folds usually 5 or 10 to be trained on each of the fold and evaluates model performance on the held-out fold to be able to see how our model would perform on different datasets through providing evaluation score such as root mean square error. We also used Randomized Search cross validation for hyperparameter tuning of models using **RandomizedSearchCV()** which select random value for each hyperparameter at every iteration which can improve our model's performance and has reduced computational costs than other

techniques. Different values were used for each hyperparameter. For model evaluation performance we used three techniques which are Mean Average Error, Root Mean Square Error, and R^2 square.

Model building steps that were applied for each model:

Model = model() – to create instance of model.

Model.fit(X_train_scaled, y_train) – to fit the model on training data.

Model_y_predict = model.predict(X_test_scaled) – to use model for prediction.

The following functions were used for each model:

- Lasso()
- KNeighborsRegressor()
- SVR()
- DecisionTreeRegressor()
- RandomForestRegressor()
- GradientBoostingRegressor()
- XGBRegressor()
- lgb.LGBMRegressor()
- VotingRegressor()

Chapter 4

4.0 Findings

This chapter contains the final result regarding the previously built models for Ames Housing dataset house price predictions using the parameters represented in the following table. The models were evaluated based on three key performance metrics which are R^2 squared, Mean Absolute Error (MAE), Root Mean Squared Error (RMSE) that are widely used in various regressions as they provide comprehensive view of model's performance.

4.1 Hyperparameter tuning results

We used the highlighted parameters in the following table which resulted from Randomized Search cross validation technique to enhance models' performance while for some models such as SVR and RFR hyperparameter tuning decreased their performance so, we used default parameters for each of them.

Tabel 4. 1 Hyperparameter Tuning Values.

Model	Hyperparameter	Values
Lasso	alpha	[0.001 ,0.01,0.1,0.000391,1,10,100,1000]
	random_state	[1,5, 10 ,24,42]
KNN	n_neighbors	[5,6, 7 ,8,9,10,11,12,13,14...20]
SVR	Default parameter	Default values
DTR	max_depth	[1,2,3,4,4,6,7,8.... 20 ,21]
	min_samples_split	[2,3,4,4,6,7,8,9, 10 ,11]
	min_samples_leaf	[1,2,3,4,4,6,7,8, 9 ,10,11]
	max_features	[20,30, 50]
RFR	Default parameter	Default values
GBR	n_estimators	[2,3,4,5,6,7,20, 30 ,50]
	alpha	[0.001,0.1, 0.5 ,0.2]
XGBoost	n_estimators	[100,200,300,500,2200, 220]
	learning_rate	[0.01,0.03, 0.05 ,0.1,0.001]
	max_depth	[3,4, 5 ,6]
	min_child_weight	[1,2,3,4, 1.7817 ,0.178,0.0178]

LGBM	gamma	[0 ,0.1,0.2,0.3,0.005,0.0045,0.35,0.5]
	colsample_bytree	[0.6, 0.7 ,0.8,0.9]
	alpha	[0,0.1, 0.5 ,1]
	random_state	42
	learning_rate	[0.7,0.8,0.1,0.001, 0.2 ,0.0003,0.3,0.5]
	max_depth	[3,4,5, 6 ,50]
	n_estimators	[6,7,10,11, 50 ,60]
	num_leaves	[5, 6 ,7]
	reg_alpha	[0 ,0.1,0.5,1]
Voting	Weights:	[0.2,0.6,0.2]
	[gbr_model,xgb_model,lgb_model]	[0.6,0.2,0.2][0.40,0.40,0.40][0.20,0.40,0.50]

4.2 Model Performance

The following table and figures summarize each model performance:

Tabel 4. 2 Model Evaluation Comparison.

Model Name	R2 Score	MAE Score	RMSE Score
Lasso	0.904482	0.089703	0.126982
Decision Tree	0.791153	0.144035	0.187764
Support Vector Regressor	0.821243	0.111036	0.173712
KNN	0.818839	0.128605	0.174876
Random Forest	0.894217	0.094702	0.133631
Gradient Boosting	0.877650	0.104677	0.143715
XGBoost	0.913757	0.085211	0.120659
LightGBM	0.910886	0.088508	0.122651
Voting Regressor	0.912203	0.086321	0.121741

The result showed that XGBoost model comes in the first place as the top performing model achieving the highest R^2 squared score (0.913757), lowest MAE score (0.085211) and lowest RMSE score (0.120659). The Second-best performing model is the voting regressor which combines Gradient boosting, XGBoost, and LightGBM models achieving high R^2 squared score (0.912203), low MAE score (0.086321) and low RMSE (0.121741) . Also, LightGBM, and Lasso showed a strong performance achieving R^2 squared score above 0.90.

while Decision Tree model considered to be the lowest performing model achieving lowest R^2 squared score (0.791153), highest MAE score (0.144035) and lowest RMSE score (0.187764). For further investigation of the most features that were used by XGBoost that contributed the most for predicting house prices, we used the following code `xgb_model.feature_importances_` our champion model, as a result we found that top 5 most important features that contributed the most in our model predictive power which are OverallQual (30%) which is overall quality of the house, ExterQual (14%) which is exterior quality, GarageCars (11%) which is size of garage in terms of number of cars, GrLivArea (5%) which is above ground area and KitchenQual (5%) which is kitchen quality.

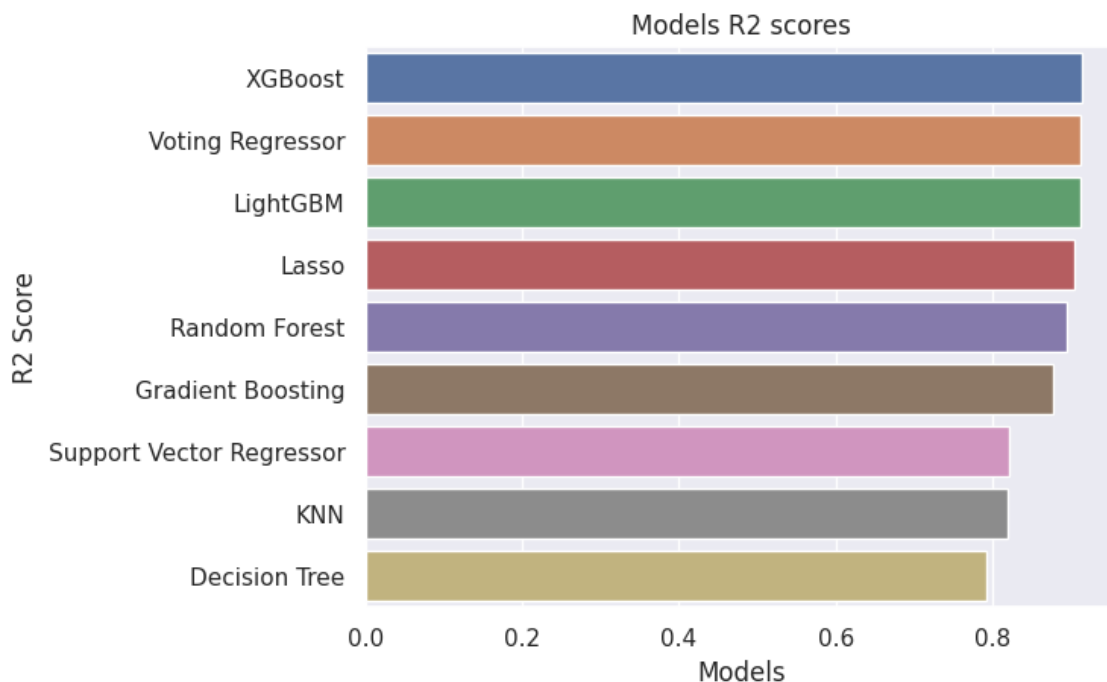


Figure 4. 1 Models R2 Scores.




Figure 4. 2 Models MAE Scores.

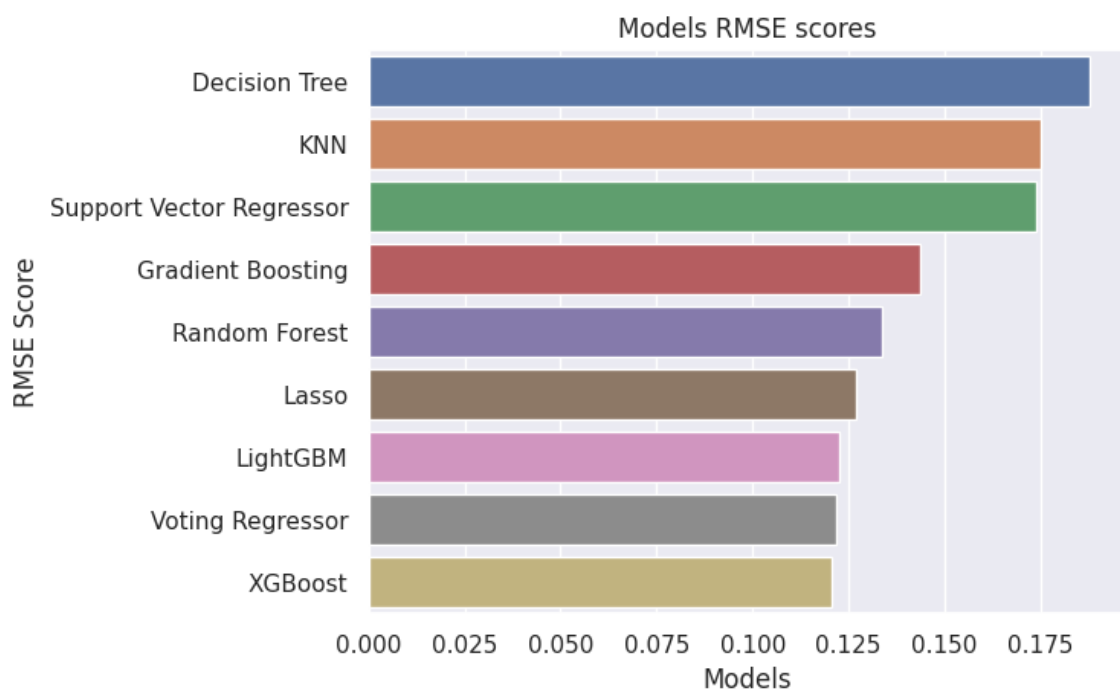


Figure 4. 3 Models RMSE Scores.

4.4 Discussion

The XGBoost is the champion model, its performance results suggest that it is a powerful ensemble model that was effective in capturing the underlying patterns within the Ames Housing dataset. Furthermore, the strong performance of other advanced models including Voting Regressor which combined between (gradient boosting, XGboost and LightGBM) models, as well as LightGBM performance. All the previous mentioned models concludes that the house price predictions can be quite complicated, and advanced regression methods can be used as they combine the advantages of various models for predictions.

Chapter 5

5.0 Conclusion

In this study we covered a lot of house price predictions aspects. We gave a brief overlook to fourth industrial revolution and property technology impact on Real Estate industry. In addition, to mentioning previous research about house price prediction. We also walked through the importance and process of house price prediction using Ames' housing dataset from Kaggle.

We retrieved the data and preformed exploratory data analysis to have a basic understand to it, then we conducted necessary preprocessing techniques including missing values imputation, feature engineering, log transformation the target variable, handling outliers, splitting, and standardizing data to prepare it in a suitable format that can be used for modelling.

Then, we built nine popular models including linear, non-linear and ensemble models for house price predictions, we performed cross validation for every model to see how the models will perform on unseen data, in addition to hyperparameter tuning using randomized search cross validation technique to enhance models' performance.

After constructing the models, we evaluated them using R^2 squared, MAE, and RMSE metrics. According to our final model evaluation XGBoost was our champion model through achieving highest R^2 , lowest MAE and RMSE followed by Voting model achieving second best performance. While the lowest performing model was Decision Tree model the achieved lowest R^2 and highest MAE and RMSE. It is important to understand that that there is no single model that can work for every problem case, the selection of appropriate model depends on several factors including data quality, and problem complexity.

Finally, Ames, Iowa city can benefit from using the XGBoost model in the process of house price predictions as it will make the process more efficient and save lots of resources including time, effort and cost.

5.1 Limitations

This study contains some limitations. First, the small size of data which only contains 1460 observations which can result in lack of capturing the diversity and variability of house market trends that could affect predictions. Furthermore, small dataset can affect models' ability to generalize well on unseen data , as well as inability to fully capture complex patterns.

Second, the data contains a large number of subjective features including but not limited to OverallQual, ExterQual, KitchenQual, and BsmtQual which are rated poor to excellent which can be exposed to bias. Moreover, the dataset lacks necessary features that may also affect the house prices such as environmental conditions, population, and proximity to desired amenities (transportation, schools, parks, and markets).

Third, the same data structure and features are necessary in order to be able to use the model built in this study for future predictions.

Fourth, as house price prediction is important, however it's not enough and needs the support of experience in the real estate field, which will be complementary to the prediction.

Lastly, that Real estate market is a dynamic industry and the dataset used in this study doesn't cover wide range of years, using a dataset that covers longer timeframe can be more effective and efficient in understanding house prices.

5.2 Recommendations

There are a number of recommendations that could be made. First, more feature engineering can be done to enhance model performance, in addition to trying different hyperparameter tuning techniques that can work better.

Second, this study can be repeated using more advanced models such as Neural Network and deep learning, which mainly wasn't applied due to small size of data.

Third, more study can be performed to enhance house price predictions process to be recognized as a successful opportunity for real-estate companies and to encourage them to use property technology.

Lastly, for the future a website or mobile application for house price predictions can be built and used in order to help buyers, investors and agencies to better understand the real estate market and not to be scammed with incorrect property prices or costs.

References

- Abdul-Rahman, S., Zulkifley, N. H., Ismail, I., & Mutalib, S. (2021). Advanced machine learning algorithms for house price prediction: Case study in Kuala Lumpur. *International Journal of Advanced Computer Science and Applications*, 12(12).
<https://search.proquest.com/openview/a0ba078716308553044e6f39d9131222/1?pq-origsite=gscholar&cbl=5444811>
- Albon, C. (2018). *Machine learning with python cookbook: Practical solutions from preprocessing to deep learning*. O'Reilly Media, Inc.
https://books.google.com/books?hl=en&lr=&id=kIhQDwAAQBAJ&oi=fnd&pg=PT107&dq=Machine+Learning+with+Python+Cookbook+:+Practical+Solutions+&ots=OoQt0KhhJO&sig=_T6fmYAUyvdCxPYZalOTFgg-P_o
- Alyami, M., Khan, M., Fawad, M., Nawaz, R., Hammad, A. W. A., Najeh, T., & Gamil, Y. (2024). Predictive modeling for compressive strength of 3D printed fiber-reinforced concrete using machine learning algorithms. *Case Studies in Construction Materials*, 20, e02728. <https://doi.org/10.1016/j.cscm.2023.e02728>
- Bai, C., Dallasega, P., Orzes, G., & Sarkis, J. (2020). Industry 4.0 technologies assessment: A sustainability perspective. *International Journal of Production Economics*, 229, 107776. <https://doi.org/10.1016/j.ijpe.2020.107776>
- Bansal, M., Goyal, A., & Choudhary, A. (2022). A comparative analysis of K-Nearest Neighbor, Genetic, Support Vector Machine, Decision Tree, and Long Short Term Memory algorithms in machine learning. *Decision Analytics Journal*, 3, 100071. <https://doi.org/10.1016/j.dajour.2022.100071>

- Baum, A., Saull, A., & Braesemann, F. (2020). PropTech 2020: The future of real estate. *Saïd Business School. University of Oxford Research. Search In.*
- Bittini, J. S., Rambaud, S. C., Pascual, J. L., & Moro-Visconti, R. (2022). Business Models and Sustainability Plans in the FinTech, InsurTech, and PropTech Industry: Evidence from Spain. *Sustainability (Switzerland)*, 14(19). Scopus. <https://doi.org/10.3390/su141912088>
- Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32.
<https://doi.org/10.1023/A:1010933404324>
- Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- City of Ames, IA | Home. (n.d.). Retrieved January 24, 2024, from <https://www.cityofames.org/home>
- cycles, T. text provides general information S. assumes no liability for the information given being complete or correct D. to varying update, & Text, S. C. D. M. up-to-D. D. T. R. in the. (n.d.). *Topic: Proptech*. Statista. Retrieved December 3, 2023, from <https://www.statista.com/topics/5230/proptech/>
- Dogaru, L. (2020). The Main Goals of the Fourth Industrial Revolution. Renewable Energy Perspectives. *Procedia Manufacturing*, 46, 397–401.
<https://doi.org/10.1016/j.promfg.2020.03.058>
- El Mouna, L., Silkan, H., Haynf, Y., Nann, M. F., & Tekouabou, S. C. K. (2023). *A Comparative Study of Urban House Price Prediction using Machine Learning Algorithms*. 418. Scopus. <https://doi.org/10.1051/e3sconf/202341803001>

- Geerts, M., vanden Broucke, S., & De Weerd, J. (2023). A Survey of Methods and Input Data Types for House Price Prediction. *ISPRS International Journal of Geo-Information*, 12(5), Article 5. <https://doi.org/10.3390/ijgi12050200>
- Géron, A. (2022). *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*. O'Reilly Media, Inc.
- GÓRSKA, A., MAZURCZAK, A., & STRĄCZKOWSKI, \Lukasz. (2022). IMPLEMENTATION OF MODERN TECHNOLOGIES (PROPTech) BY DEVELOPERS ON THE LOCAL HOUSING MARKET. *Scientific Papers of Silesian University of Technology. Organization & Management/Zeszyty Naukowe Politechniki Śląskiej. Seria Organizacji i Zarządzanie*, 162. <https://managementpapers.polsl.pl/wp-content/uploads/2022/12/162-G%C3%B3rska-Mazurczak-Str%C4%85czkowski-2.pdf>
- Hanuma Reddy, G., & Sriramya, P. (2022). Real Estate Price Prediction and Analysis Using Voting Regression Compared with Linear Regression. In *Advances in Parallel Computing Algorithms, Tools and Paradigms* (pp. 625–630). IOS Press. <https://doi.org/10.3233/APC220089>
- Ho, W. K. O., Tang, B.-S., & Wong, S. W. (2021). Predicting property prices with machine learning algorithms. *Journal of Property Research*, 38(1), 48–70. <https://doi.org/10.1080/09599916.2020.1832558>
- Kansal, M., Singh, P., Shukla, S., & Srivastava, S. (2023). A Comparative Study of Machine Learning Models for House Price Prediction and Analysis in Smart Cities. In F. Ortiz-Rodríguez, S. Tiwari, P. Usoro Usip, & R. Palma (Eds.), *Electronic Governance with Emerging Technologies* (pp. 168–184). Springer Nature Switzerland. https://doi.org/10.1007/978-3-031-43940-7_14

- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017).
LightGBM: A highly efficient gradient boosting decision tree. *Proceedings of the 31st International Conference on Neural Information Processing Systems*, 3149–3157.
- Kumar, S., & Syed, M. H. (2023). House Pricing Prediction Based on Composite Facility Score Using Machine Learning Algorithms. In P. K. Singh, S. T. Wierzchoń, S. Tanwar, J. J. P. C. Rodrigues, & M. Ganzha (Eds.), *Proceedings of Third International Conference on Computing, Communications, and Cyber-Security* (pp. 235–248). Springer Nature.
https://doi.org/10.1007/978-981-19-1142-2_18
- Machkour, B., & Abriane, A. (2020). Industry 4.0 and its Implications for the Financial Sector. *Procedia Computer Science*, 177, 496–502. <https://doi.org/10.1016/j.procs.2020.10.068>
- pandas documentation—Pandas 2.1.4 documentation*. (n.d.). Retrieved January 8, 2024, from <https://pandas.pydata.org/docs/index.html>
- Putatunda, S. (2019). *PropTech for Proactive Pricing of Houses in Classified Advertisements in the Indian Real Estate Market* (arXiv:1904.05328). arXiv.
<https://doi.org/10.48550/arXiv.1904.05328>
- Research, Z. M. (n.d.). *PropTech Market Size, Share, Trends, Demand, Applications 2023-2030*. Zion Market Research. Retrieved January 11, 2024, from <https://www.zionmarketresearch.com/report/prop-tech-market>
- Saiz, A. (2020). Bricks, mortar, and proptech: The economics of IT in brokerage, space utilization and commercial real estate finance. *Journal of Property Investment & Finance*, 38(4), 327–347. <https://doi.org/10.1108/JPIF-10-2019-0139>

- Sibindi, R., Mwangi, R. W., & Waititu, A. G. (2023). A boosting ensemble learning based hybrid light gradient boosting machine and extreme gradient boosting model for predicting house prices. *Engineering Reports*, 5(4), e12599. <https://doi.org/10.1002/eng2.12599>
- Starr, C. W., Saginor, J., & Worzala, E. (2020). The rise of PropTech: Emerging industrial technologies and their impact on real estate. *Journal of Property Investment & Finance*, 39(2), 157–169. <https://doi.org/10.1108/JPIF-08-2020-0090>
- Truong, Q., Nguyen, M., Dang, H., & Mei, B. (2020). Housing price prediction via improved machine learning techniques. *Procedia Computer Science*, 174, 433–442.
- Vemuri, V. K. (2020). The Hundred-Page Machine Learning Book: By Andriy Burkov, Quebec City, Canada, 2019, 160 pp., \$49.99 (Hardcover); \$29.00 (paperback); \$25.43 (Kindle Edition), (Alternatively, can purchase at leanpub.com at a minimum price of \$20.00), ISBN 978-1999579517. *Journal of Information Technology Case and Application Research*, 22(2), 136–138. <https://doi.org/10.1080/15228053.2020.1766224>
- Xu, M., David, J. M., & Kim, S. H. (2018). The fourth industrial revolution: Opportunities and challenges. *International Journal of Financial Research*, 9(2), 90–95.
- Yadav, S., Dhanda, N., Sahai, A., Verma, R., & Pandey, S. (2023). Real Estate Price Prediction Using Machine Learning. In D. K. Sarkar, P. K. Sadhu, S. Bhunia, J. Samanta, & S. Paul (Eds.), *Proceedings of the 4th International Conference on Communication, Devices and Computing* (pp. 95–111). Springer Nature. https://doi.org/10.1007/978-981-99-2710-4_9
- Zhou, J., Hai, T., Maxwell-Chigozie, E. C., Adedayo, A., Chen, Y., Iwendi, C., & Boulouard, Z. (2023). Effective House Price Prediction Using Machine Learning. In C. Iwendi, Z. Boulouard, & N. Kryvinska (Eds.), *Proceedings of ICACTCE'23—The International Conference on Advances in Communication Technology and Computer Engineering* (Vol.

735, pp. 425–436). Springer Nature Switzerland. https://doi.org/10.1007/978-3-031-37164-6_32

Zillow: Real Estate, Apartments, Mortgages & Home Values. (n.d.). Retrieved December 3, 2023, from <https://www.zillow.com/>

Appendix

Libraries Imported

```
#Import necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.feature_selection import RFE
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import Lasso, LinearRegression
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.model_selection import RandomizedSearchCV
import lightgbm as lgb
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.ensemble import VotingRegressor
```

Data Description

```
# Open and print data description file to understand data.
with open('/content/data_description.txt', 'r') as file:
    data_description = file.read()
print(data_description)
```

MSSubClass: Identifies the type of dwelling involved in the sale.

20	1-STORY 1946 & NEWER ALL STYLES
30	1-STORY 1945 & OLDER
40	1-STORY W/FINISHED ATTIC ALL AGES
45	1-1/2 STORY - UNFINISHED ALL AGES
50	1-1/2 STORY FINISHED ALL AGES
60	2-STORY 1946 & NEWER
70	2-STORY 1945 & OLDER
75	2-1/2 STORY ALL AGES
80	SPLIT OR MULTI-LEVEL
85	SPLIT FOYER
90	DUPLEX - ALL STYLES AND AGES
120	1-STORY PUD (Planned Unit Development) - 1946 & NEWER
150	1-1/2 STORY PUD - ALL AGES
160	2-STORY PUD - 1946 & NEWER
180	PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
190	2 FAMILY CONVERSION - ALL STYLES AND AGES

MSZoning: Identifies the general zoning classification of the sale.

A	Agriculture
C	Commercial
FV	Floating Village Residential
I	Industrial
RH	Residential High Density
RL	Residential Low Density
RP	Residential Low Density Park

RM Residential Medium Density

LotFrontage: Linear feet of street connected to property

LotArea: Lot size in square feet

Street: Type of road access to property

Grv1	Gravel
Pave	Paved

Alley: Type of alley access to property

Grv1	Gravel
Pave	Paved
NA	No alley access

LotShape: General shape of property

Reg	Regular
IR1	Slightly irregular
IR2	Moderately Irregular
IR3	Irregular

LandContour: Flatness of the property

Lvl	Near Flat/Level
Bnk	Banked - Quick and significant rise from street grade to building
HLS	Hillside - Significant slope from side to side
Low	Depression

Utilities: Type of utilities available

AllPub	All public Utilities (E,G,W,& S)
NoSewr	Electricity, Gas, and Water (Septic Tank)
NoSeWa	Electricity and Gas Only
ELO	Electricity only

LotConfig: Lot configuration

Inside	Inside lot
Corner	Corner lot
CulDSac	Cul-de-sac
FR2	Frontage on 2 sides of property
FR3	Frontage on 3 sides of property

LandSlope: Slope of property

Gtl	Gentle slope
Mod	Moderate Slope
Sev	Severe Slope

Neighborhood: Physical locations within Ames city limits

Blmngtn	Bloomington Heights
Blueste	Bluestem
BrDale	Briardale

BrkSide Brookside
 ClearCr Clear Creek
 CollgCr College Creek
 Crawfor Crawford
 Edwards Edwards
 Gilbert Gilbert
 IDOTRR Iowa DOT and Rail Road
 MeadowV Meadow Village
 Mitchel Mitchell
 Names North Ames
 NoRidge Northridge
 NPKVill Northpark Villa
 NridgHt Northridge Heights
 NWAmes Northwest Ames
 OldTown Old Town
 SWISU South & West of Iowa State University
 Sawyer Sawyer
 SawyerW Sawyer West
 Somerst Somerset
 StoneBr Stone Brook
 Timber Timberland
 Veenker Veenker

Condition1: Proximity to various conditions

Artery Adjacent to arterial street
 Feedr Adjacent to feeder street
 Norm Normal
 RRNn Within 200' of North-South Railroad
 RRAn Adjacent to North-South Railroad
 PosN Near positive off-site feature--park, greenbelt, etc.
 PosA Adjacent to postive off-site feature
 RRNe Within 200' of East-West Railroad
 RRAe Adjacent to East-West Railroad

Condition2: Proximity to various conditions (if more than one is present)

Artery Adjacent to arterial street
 Feedr Adjacent to feeder street
 Norm Normal
 RRNn Within 200' of North-South Railroad
 RRAn Adjacent to North-South Railroad
 PosN Near positive off-site feature--park, greenbelt, etc.
 PosA Adjacent to postive off-site feature
 RRNe Within 200' of East-West Railroad
 RRAe Adjacent to East-West Railroad

BldgType: Type of dwelling

1Fam Single-family Detached
 2FmCon Two-family Conversion; originally built as one-family dwelling
 Duplx Duplex
 TwnhsE Townhouse End Unit
 TwnhsI Townhouse Inside Unit

HouseStyle: Style of dwelling

1Story One story
 1.5Fin One and one-half story: 2nd level finished

1.5Unf One and one-half story: 2nd level unfinished
 2Story Two story
 2.5Fin Two and one-half story: 2nd level finished
 2.5Unf Two and one-half story: 2nd level unfinished
 SFoyer Split Foyer
 SLvl Split Level

OverallQual: Rates the overall material and finish of the house

10 Very Excellent
 9 Excellent
 8 Very Good
 7 Good
 6 Above Average
 5 Average
 4 Below Average
 3 Fair
 2 Poor
 1 Very Poor

OverallCond: Rates the overall condition of the house

10 Very Excellent
 9 Excellent
 8 Very Good
 7 Good
 6 Above Average
 5 Average
 4 Below Average
 3 Fair
 2 Poor
 1 Very Poor

YearBuilt: Original construction date

YearRemodAdd: Remodel date (same as construction date if no remodeling or additions)

RoofStyle: Type of roof

Flat Flat
 Gable Gable
 Gambrel Gambrel (Barn)
 Hip Hip
 Mansard Mansard
 Shed Shed

RoofMatl: Roof material

ClyTile Clay or Tile
 CompShg Standard (Composite) Shingle
 Membran Membrane
 Metal Metal
 Roll Roll
 Tar&Grv Gravel & Tar
 WdShake Wood Shakes
 WdShngl Wood Shingles

Exterior1st: Exterior covering on house

AsbShng Asbestos Shingles
 AsphShn Asphalt Shingles
 BrkComm Brick Common
 BrkFace Brick Face
 CBlock Cinder Block
 CemntBd Cement Board
 HdBoard Hard Board
 ImStucc Imitation Stucco
 MetalSd Metal Siding
 Other Other
 Plywood Plywood
 PreCast PreCast
 Stone Stone
 Stucco Stucco
 VinylSd Vinyl Siding
 Wd Sdng Wood Siding
 WdShng Wood Shingles

Exterior2nd: Exterior covering on house (if more than one material)

AsbShng Asbestos Shingles
 AsphShn Asphalt Shingles
 BrkComm Brick Common
 BrkFace Brick Face
 CBlock Cinder Block
 CemntBd Cement Board
 HdBoard Hard Board
 ImStucc Imitation Stucco
 MetalSd Metal Siding
 Other Other
 Plywood Plywood
 PreCast PreCast
 Stone Stone
 Stucco Stucco
 VinylSd Vinyl Siding
 Wd Sdng Wood Siding
 WdShng Wood Shingles

MasVnrType: Masonry veneer type

BrkCmn Brick Common
 BrkFace Brick Face
 CBlock Cinder Block
 None None
 Stone Stone

MasVnrArea: Masonry veneer area in square feet

ExterQual: Evaluates the quality of the material on the exterior

Ex Excellent
 Gd Good
 TA Average/Typical
 Fa Fair
 Po Poor

ExterCond: Evaluates the present condition of the material on the exterior

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
Po	Poor

Foundation: Type of foundation

BrkTil	Brick & Tile
CBlock	Cinder Block
PConc	Poured Contrete
Slab	Slab
Stone	Stone
Wood	Wood

BsmtQual: Evaluates the height of the basement

Ex	Excellent (100+ inches)
Gd	Good (90-99 inches)
TA	Typical (80-89 inches)
Fa	Fair (70-79 inches)
Po	Poor (<70 inches)
NA	No Basement

BsmtCond: Evaluates the general condition of the basement

Ex	Excellent
Gd	Good
TA	Typical - slight dampness allowed
Fa	Fair - dampness or some cracking or settling
Po	Poor - Severe cracking, settling, or wetness
NA	No Basement

BsmtExposure: Refers to walkout or garden level walls

Gd	Good Exposure
Av	Average Exposure (split levels or foyers typically score average or above)
Mn	Mimimum Exposure
No	No Exposure
NA	No Basement

BsmtFinType1: Rating of basement finished area

GLQ	Good Living Quarters
ALQ	Average Living Quarters
BLQ	Below Average Living Quarters
Rec	Average Rec Room
LwQ	Low Quality
Unf	Unfinished
NA	No Basement

BsmtFinSF1: Type 1 finished square feet

BsmtFinType2: Rating of basement finished area (if multiple types)

GLQ	Good Living Quarters
ALQ	Average Living Quarters
BLQ	Below Average Living Quarters

Rec	Average Rec Room
LwQ	Low Quality
Unf	Unfinished
NA	No Basement

BsmtFinSF2: Type 2 finished square feet

BsmtUnfSF: Unfinished square feet of basement area

TotalBsmtSF: Total square feet of basement area

Heating: Type of heating

Floor	Floor Furnace
GasA	Gas forced warm air furnace
GasW	Gas hot water or steam heat
Grav	Gravity furnace
OthW	Hot water or steam heat other than gas
Wall	Wall furnace

HeatingQC: Heating quality and condition

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
Po	Poor

CentralAir: Central air conditioning

N	No
Y	Yes

Electrical: Electrical system

SBrkr	Standard Circuit Breakers & Romex
FuseA	Fuse Box over 60 AMP and all Romex wiring (Average)
FuseF	60 AMP Fuse Box and mostly Romex wiring (Fair)
FuseP	60 AMP Fuse Box and mostly knob & tube wiring (poor)
Mix	Mixed

1stFlrSF: First Floor square feet

2ndFlrSF: Second floor square feet

LowQualFinSF: Low quality finished square feet (all floors)

GrLivArea: Above grade (ground) living area square feet

BsmtFullBath: Basement full bathrooms

BsmtHalfBath: Basement half bathrooms

FullBath: Full bathrooms above grade

HalfBath: Half baths above grade

Bedroom: Bedrooms above grade (does NOT include basement bedrooms)

Kitchen: Kitchens above grade

KitchenQual: Kitchen quality

Ex	Excellent
Gd	Good
TA	Typical/Average
Fa	Fair
Po	Poor

TotRmsAbvGrd: Total rooms above grade (does not include bathrooms)

Functional: Home functionality (Assume typical unless deductions are warranted)

Typ	Typical Functionality
Min1	Minor Deductions 1
Min2	Minor Deductions 2
Mod	Moderate Deductions
Maj1	Major Deductions 1
Maj2	Major Deductions 2
Sev	Severely Damaged
Sal	Salvage only

Fireplaces: Number of fireplaces

FireplaceQu: Fireplace quality

Ex	Excellent - Exceptional Masonry Fireplace
Gd	Good - Masonry Fireplace in main level
TA	Average - Prefabricated Fireplace in main living area or Masonry Fireplace in basement
Fa	Fair - Prefabricated Fireplace in basement
Po	Poor - Ben Franklin Stove
NA	No Fireplace

GarageType: Garage location

2Types	More than one type of garage
Attchd	Attached to home
Basment	Basement Garage
BuiltIn	Built-In (Garage part of house - typically has room above garage)
CarPort	Car Port
Detchd	Detached from home
NA	No Garage

GarageYrBlt: Year garage was built

GarageFinish: Interior finish of the garage

Fin	Finished
RFn	Rough Finished
Unf	Unfinished
NA	No Garage

GarageCars: Size of garage in car capacity

GarageArea: Size of garage in square feet

GarageQual: Garage quality

Ex	Excellent
Gd	Good
TA	Typical/Average
Fa	Fair
Po	Poor
NA	No Garage

GarageCond: Garage condition

Ex	Excellent
Gd	Good
TA	Typical/Average
Fa	Fair
Po	Poor
NA	No Garage

PavedDrive: Paved driveway

Y	Paved
P	Partial Pavement
N	Dirt/Gravel

WoodDeckSF: Wood deck area in square feet

OpenPorchSF: Open porch area in square feet

EnclosedPorch: Enclosed porch area in square feet

3SsnPorch: Three season porch area in square feet

ScreenPorch: Screen porch area in square feet

PoolArea: Pool area in square feet

PoolQC: Pool quality

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
NA	No Pool

Fence: Fence quality

GdPrv	Good Privacy
MnPrv	Minimum Privacy
GdWo	Good Wood
MnWw	Minimum Wood/Wire
NA	No Fence

MiscFeature: Miscellaneous feature not covered in other categories

Elev	Elevator
Gar2	2nd Garage (if not described in garage section)
Othr	Other
Shed	Shed (over 100 SF)
TenC	Tennis Court
NA	None

MiscVal: \$Value of miscellaneous feature

MoSold: Month Sold (MM)

YrSold: Year Sold (YYYY)

SaleType: Type of sale

WD	Warranty Deed - Conventional
CWD	Warranty Deed - Cash
VWD	Warranty Deed - VA Loan
New	Home just constructed and sold
COD	Court Officer Deed/Estate
Con	Contract 15% Down payment regular terms
ConLw	Contract Low Down payment and low interest
ConLI	Contract Low Interest
ConLD	Contract Low Down
Oth	Other

SaleCondition: Condition of sale

Normal	Normal Sale
Abnorml	Abnormal Sale - trade, foreclosure, short sale
AdjLand	Adjoining Land Purchase
Alloca	Allocation - two linked properties with separate deeds, typically condo with a garage unit
Family	Sale between family members
Partial	Home was not completed when last assessed (associated with New Homes)

###Data Acquisiton

#upload the training data.

```
house_df = pd.read_csv('/content/train.csv')
```

###Data Exploration

#Check head of data

```
house_df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	

```
4          Lvl    AllPub    ...          0    NaN    NaN          NaN          0    12
```

```
YrSold SaleType SaleCondition SalePrice
0    2008         WD         Normal    208500
1    2007         WD         Normal    181500
2    2008         WD         Normal    223500
3    2006         WD        Abnorml    140000
4    2008         WD         Normal    250000
```

```
[5 rows x 81 columns]
```

```
house_df.columns
```

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
       'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
       'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
       'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
       'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
       'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
       'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
       'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
       'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
       'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
       'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
       'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
       'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
       'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
       'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
       'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
       'SaleCondition', 'SalePrice'],
      dtype='object')
```

```
# Understand the shape of training data.
```

```
house_df.shape
```

```
(1460, 81)
```

```
# check for duplicates
```

```
house_df.duplicated().sum()
```

```
0
```

```
# Drop Id column
```

```
house_df.drop(['Id'], axis = 1, inplace = True)
```

```
house_df.head()
```

```
MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0          60      RL         65.0      8450  Pave   NaN      Reg
1          20      RL         80.0      9600  Pave   NaN      Reg
2          60      RL         68.0     11250  Pave   NaN     IR1
3          70      RL         60.0      9550  Pave   NaN     IR1
4          60      RL         84.0     14260  Pave   NaN     IR1
```

```
LandContour Utilities LotConfig ... PoolArea PoolQC Fence MiscFeature \
0          Lvl    AllPub    Inside ...          0    NaN    NaN      NaN
1          Lvl    AllPub    FR2    ...          0    NaN    NaN      NaN
2          Lvl    AllPub    Inside ...          0    NaN    NaN      NaN
3          Lvl    AllPub    Corner ...          0    NaN    NaN      NaN
4          Lvl    AllPub    FR2    ...          0    NaN    NaN      NaN
```

```
MiscVal MoSold YrSold SaleType SaleCondition SalePrice
```

0	0	2	2008	WD	Normal	208500
1	0	5	2007	WD	Normal	181500
2	0	9	2008	WD	Normal	223500
3	0	2	2006	WD	Abnorml	140000
4	0	12	2008	WD	Normal	250000

[5 rows x 80 columns]

house training data info.

house_df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1460 entries, 0 to 1459

Data columns (total 80 columns):

#	Column	Non-Null Count	Dtype
0	MSSubClass	1460 non-null	int64
1	MSZoning	1460 non-null	object
2	LotFrontage	1201 non-null	float64
3	LotArea	1460 non-null	int64
4	Street	1460 non-null	object
5	Alley	91 non-null	object
6	LotShape	1460 non-null	object
7	LandContour	1460 non-null	object
8	Utilities	1460 non-null	object
9	LotConfig	1460 non-null	object
10	LandSlope	1460 non-null	object
11	Neighborhood	1460 non-null	object
12	Condition1	1460 non-null	object
13	Condition2	1460 non-null	object
14	BldgType	1460 non-null	object
15	HouseStyle	1460 non-null	object
16	OverallQual	1460 non-null	int64
17	OverallCond	1460 non-null	int64
18	YearBuilt	1460 non-null	int64
19	YearRemodAdd	1460 non-null	int64
20	RoofStyle	1460 non-null	object
21	RoofMatl	1460 non-null	object
22	Exterior1st	1460 non-null	object
23	Exterior2nd	1460 non-null	object
24	MasVnrType	1452 non-null	object
25	MasVnrArea	1452 non-null	float64
26	ExterQual	1460 non-null	object
27	ExterCond	1460 non-null	object
28	Foundation	1460 non-null	object
29	BsmtQual	1423 non-null	object
30	BsmtCond	1423 non-null	object
31	BsmtExposure	1422 non-null	object
32	BsmtFinType1	1423 non-null	object
33	BsmtFinSF1	1460 non-null	int64
34	BsmtFinType2	1422 non-null	object
35	BsmtFinSF2	1460 non-null	int64
36	BsmtUnfSF	1460 non-null	int64
37	TotalBsmtSF	1460 non-null	int64
38	Heating	1460 non-null	object
39	HeatingQC	1460 non-null	object
40	CentralAir	1460 non-null	object
41	Electrical	1459 non-null	object
42	1stFlrSF	1460 non-null	int64
43	2ndFlrSF	1460 non-null	int64
44	LowQualFinSF	1460 non-null	int64


```

45 GrLivArea      1460 non-null int64
46 BsmtFullBath  1460 non-null int64
47 BsmtHalfBath  1460 non-null int64
48 FullBath      1460 non-null int64
49 HalfBath      1460 non-null int64
50 BedroomAbvGr  1460 non-null int64
51 KitchenAbvGr  1460 non-null int64
52 KitchenQual    1460 non-null object
53 TotRmsAbvGrd   1460 non-null int64
54 Functional     1460 non-null object
55 Fireplaces     1460 non-null int64
56 FireplaceQu    770 non-null object
57 GarageType     1379 non-null object
58 GarageYrBlt    1379 non-null float64
59 GarageFinish   1379 non-null object
60 GarageCars     1460 non-null int64
61 GarageArea     1460 non-null int64
62 GarageQual     1379 non-null object
63 GarageCond     1379 non-null object
64 PavedDrive     1460 non-null object
65 WoodDeckSF     1460 non-null int64
66 OpenPorchSF    1460 non-null int64
67 EnclosedPorch  1460 non-null int64
68 3SsnPorch      1460 non-null int64
69 ScreenPorch    1460 non-null int64
70 PoolArea       1460 non-null int64
71 PoolQC         7 non-null object
72 Fence          281 non-null object
73 MiscFeature    54 non-null object
74 MiscVal        1460 non-null int64
75 MoSold         1460 non-null int64
76 YrSold         1460 non-null int64
77 SaleType       1460 non-null object
78 SaleCondition  1460 non-null object
79 SalePrice      1460 non-null int64
dtypes: float64(3), int64(34), object(43)
memory usage: 912.6+ KB

```

```

# Describe the training data.
house_df.describe().T

```

	count	mean	std	min	25%	\
MSSubClass	1460.0	56.897260	42.300571	20.0	20.00	
LotFrontage	1201.0	70.049958	24.284752	21.0	59.00	
LotArea	1460.0	10516.828082	9981.264932	1300.0	7553.50	
OverallQual	1460.0	6.099315	1.382997	1.0	5.00	
OverallCond	1460.0	5.575342	1.112799	1.0	5.00	
YearBuilt	1460.0	1971.267808	30.202904	1872.0	1954.00	
YearRemodAdd	1460.0	1984.865753	20.645407	1950.0	1967.00	
MasVnrArea	1452.0	103.685262	181.066207	0.0	0.00	
BsmtFinSF1	1460.0	443.639726	456.098091	0.0	0.00	
BsmtFinSF2	1460.0	46.549315	161.319273	0.0	0.00	
BsmtUnfSF	1460.0	567.240411	441.866955	0.0	223.00	
TotalBsmtSF	1460.0	1057.429452	438.705324	0.0	795.75	
1stFlrSF	1460.0	1162.626712	386.587738	334.0	882.00	
2ndFlrSF	1460.0	346.992466	436.528436	0.0	0.00	
LowQualFinSF	1460.0	5.844521	48.623081	0.0	0.00	
GrLivArea	1460.0	1515.463699	525.480383	334.0	1129.50	
BsmtFullBath	1460.0	0.425342	0.518911	0.0	0.00	
BsmtHalfBath	1460.0	0.057534	0.238753	0.0	0.00	
FullBath	1460.0	1.565068	0.550916	0.0	1.00	

HalfBath	1460.0	0.382877	0.502885	0.0	0.00
BedroomAbvGr	1460.0	2.866438	0.815778	0.0	2.00
KitchenAbvGr	1460.0	1.046575	0.220338	0.0	1.00
TotRmsAbvGrd	1460.0	6.517808	1.625393	2.0	5.00
Fireplaces	1460.0	0.613014	0.644666	0.0	0.00
GarageYrBlt	1379.0	1978.506164	24.689725	1900.0	1961.00
GarageCars	1460.0	1.767123	0.747315	0.0	1.00
GarageArea	1460.0	472.980137	213.804841	0.0	334.50
WoodDeckSF	1460.0	94.244521	125.338794	0.0	0.00
OpenPorchSF	1460.0	46.660274	66.256028	0.0	0.00
EnclosedPorch	1460.0	21.954110	61.119149	0.0	0.00
3SsnPorch	1460.0	3.409589	29.317331	0.0	0.00
ScreenPorch	1460.0	15.060959	55.757415	0.0	0.00
PoolArea	1460.0	2.758904	40.177307	0.0	0.00
MiscVal	1460.0	43.489041	496.123024	0.0	0.00
MoSold	1460.0	6.321918	2.703626	1.0	5.00
YrSold	1460.0	2007.815753	1.328095	2006.0	2007.00
SalePrice	1460.0	180921.195890	79442.502883	34900.0	129975.00

	50%	75%	max
MSSubClass	50.0	70.00	190.0
LotFrontage	69.0	80.00	313.0
LotArea	9478.5	11601.50	215245.0
OverallQual	6.0	7.00	10.0
OverallCond	5.0	6.00	9.0
YearBuilt	1973.0	2000.00	2010.0
YearRemodAdd	1994.0	2004.00	2010.0
MasVnrArea	0.0	166.00	1600.0
BsmtFinSF1	383.5	712.25	5644.0
BsmtFinSF2	0.0	0.00	1474.0
BsmtUnfSF	477.5	808.00	2336.0
TotalBsmtSF	991.5	1298.25	6110.0
1stFlrSF	1087.0	1391.25	4692.0
2ndFlrSF	0.0	728.00	2065.0
LowQualFinSF	0.0	0.00	572.0
GrLivArea	1464.0	1776.75	5642.0
BsmtFullBath	0.0	1.00	3.0
BsmtHalfBath	0.0	0.00	2.0
FullBath	2.0	2.00	3.0
HalfBath	0.0	1.00	2.0
BedroomAbvGr	3.0	3.00	8.0
KitchenAbvGr	1.0	1.00	3.0
TotRmsAbvGrd	6.0	7.00	14.0
Fireplaces	1.0	1.00	3.0
GarageYrBlt	1980.0	2002.00	2010.0
GarageCars	2.0	2.00	4.0
GarageArea	480.0	576.00	1418.0
WoodDeckSF	0.0	168.00	857.0
OpenPorchSF	25.0	68.00	547.0
EnclosedPorch	0.0	0.00	552.0
3SsnPorch	0.0	0.00	508.0
ScreenPorch	0.0	0.00	480.0
PoolArea	0.0	0.00	738.0
MiscVal	0.0	0.00	15500.0
MoSold	6.0	8.00	12.0
YrSold	2008.0	2009.00	2010.0
SalePrice	163000.0	214000.00	755000.0

```
# describe target column 'SalePrice'.
house_df['SalePrice'].describe()
```

```

count      1460.000000
mean      180921.195890
std       79442.502883
min       34900.000000
25%      129975.000000
50%      163000.000000
75%      214000.000000
max       755000.000000
Name: SalePrice, dtype: float64


```

```

# Distribution of target column 'SalePrice'.
sns.histplot(house_df['SalePrice'])
plt.title("Distribution of SalePrice")

```

```
Text(0.5, 1.0, 'Distribution of SalePrice')
```



```

# Check variance of SalePrice
house_df['SalePrice'].var()

```

```
6311111264.297448
```

```

# Separate numerical and categorical features.
numerical_house_df = house_df.select_dtypes(include=['int64', 'float64'])
categorical_house_df = house_df.select_dtypes(exclude=['int64', 'float64'])

```

```
numerical_house_df.columns
```

```

Index(['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
       'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2',

```

```

        'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
        'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
        'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
        'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
        'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',
        'MoSold', 'YrSold', 'SalePrice'],
        dtype='object')

categorical_house_df.head(3)

  MSZoning Street Alley LotShape LandContour Utilities LotConfig LandSlope \
0      RL   Pave   NaN      Reg          Lvl   AllPub   Inside    Gtl
1      RL   Pave   NaN      Reg          Lvl   AllPub   FR2      Gtl
2      RL   Pave   NaN      IR1          Lvl   AllPub   Inside    Gtl









  Neighborhood Condition1 ... GarageType GarageFinish GarageQual GarageCond \
0      CollgCr      Norm ...   Attchd          RFn          TA          TA
1      Veenker      Feedr ...   Attchd          RFn          TA          TA
2      CollgCr      Norm ...   Attchd          RFn          TA          TA

  PavedDrive PoolQC Fence MiscFeature SaleType SaleCondition
0          Y   NaN   NaN          NaN      WD      Normal
1          Y   NaN   NaN          NaN      WD      Normal
2          Y   NaN   NaN          NaN      WD      Normal

[3 rows x 43 columns]

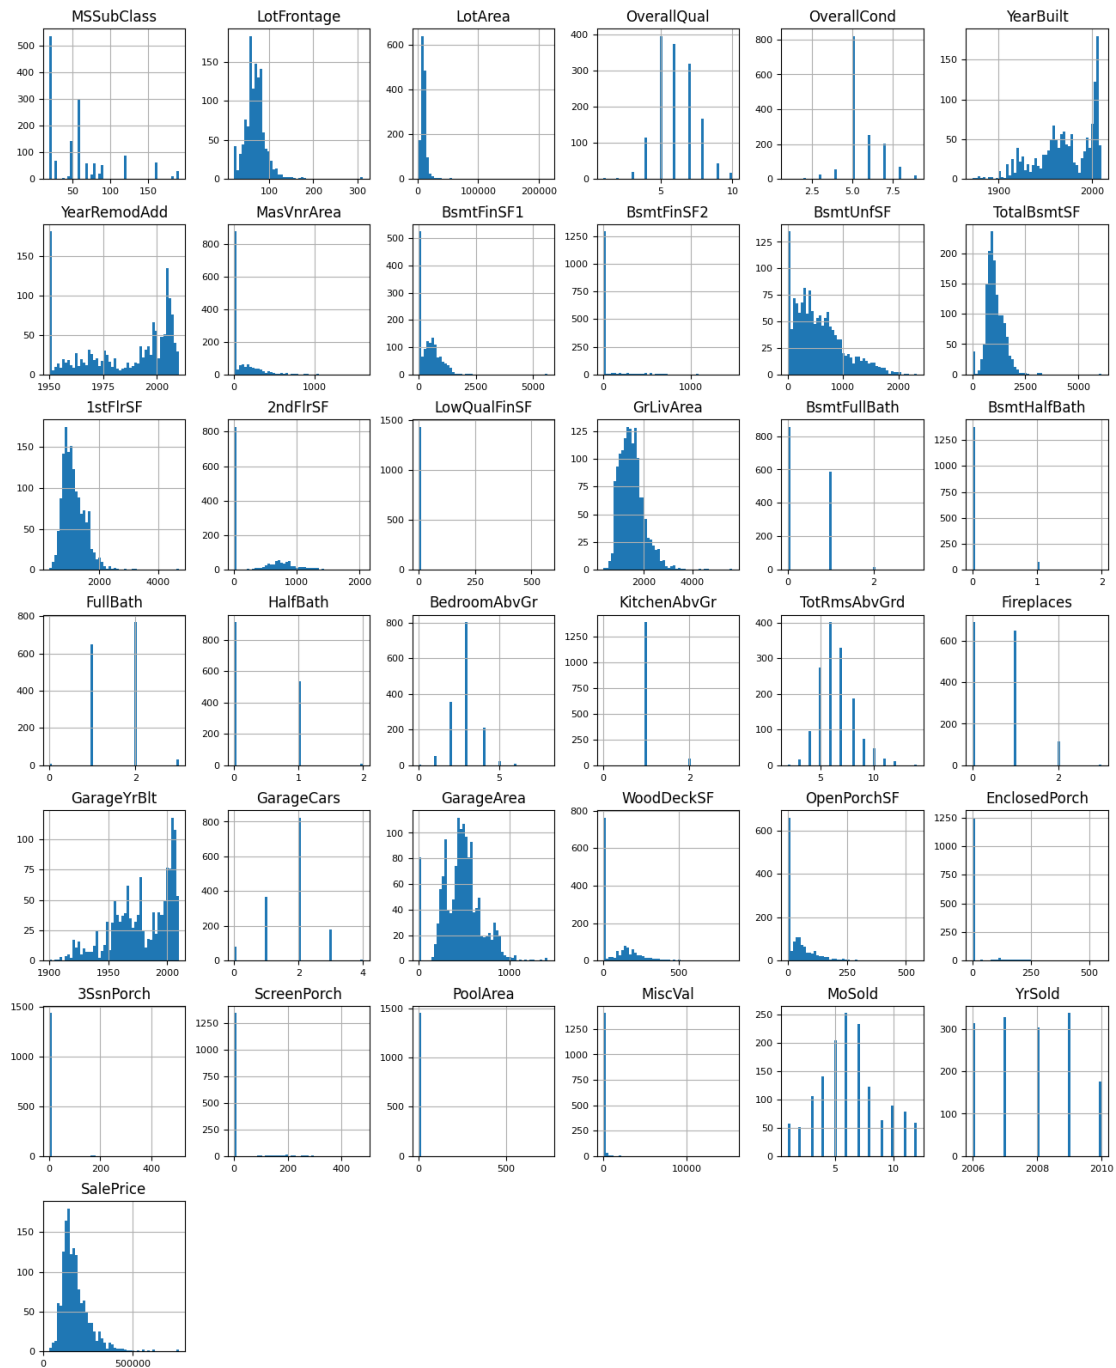
# Bar plot for exploring categorical features
cat_features = len(categorical_house_df.columns)
# Calculate the figure size
fig, axes = plt.subplots(nrows= 9, ncols= 5, figsize=(15, 40))
# Flatten the axes array to simplify indexing
axes = axes.flatten()
# Iterate over categorical features and create individual bar plot in subplots
for i, column in enumerate(categorical_house_df.columns):
    if i < cat_features:
        sns.barplot(x=house_df['SalePrice'],y=column ,data= categorical_house_df,orient = 'h',
ax=axes[i])
        axes[i].set_title(f'{column} bar plot')
# Minimize to only 43 plot not 45
for i in range(len(categorical_house_df.columns), len(axes)):
    fig.delaxes(axes[i])
#prevent overlapping title.
plt.tight_layout()
plt.show()

```




Numerical features distribution.

```
numerical_house_df.hist(figsize=(16,20), bins=50, xlabelsize=8, ylabelsize=8)
plt.title('Distribution of Numerical features')
plt.show()
```



```
# Heatmap of numerical features.
corrmat = numerical_house_df.corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, square=True, cmap='Blues')
```

<Axes: >



```
corr_matrix = numerical_house_df.corr()
corr_matrix['SalePrice'].sort_values(ascending=False)
```

```
SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmtSF    0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
YearBuilt      0.522897
YearRemodAdd    0.507101
GarageYrBlt    0.486362
MasVnrArea     0.477493
Fireplaces     0.466929
BsmtFinSF1     0.386420
```

```

LotFrontage      0.351799
WoodDeckSF       0.324413
2ndFlrSF         0.319334
OpenPorchSF      0.315856
HalfBath          0.284108
LotArea          0.263843
BsmtFullBath     0.227122
BsmtUnfSF        0.214479
BedroomAbvGr     0.168213
ScreenPorch      0.111447
PoolArea         0.092404
MoSold           0.046432
3SsnPorch        0.044584
BsmtFinSF2       -0.011378
BsmtHalfBath     -0.016844
MiscVal          -0.021190
LowQualFinSF     -0.025606
YrSold           -0.028923
OverallCond      -0.077856
MSSubClass       -0.084284
EnclosedPorch    -0.128578
KitchenAbvGr     -0.135907
Name: SalePrice, dtype: float64

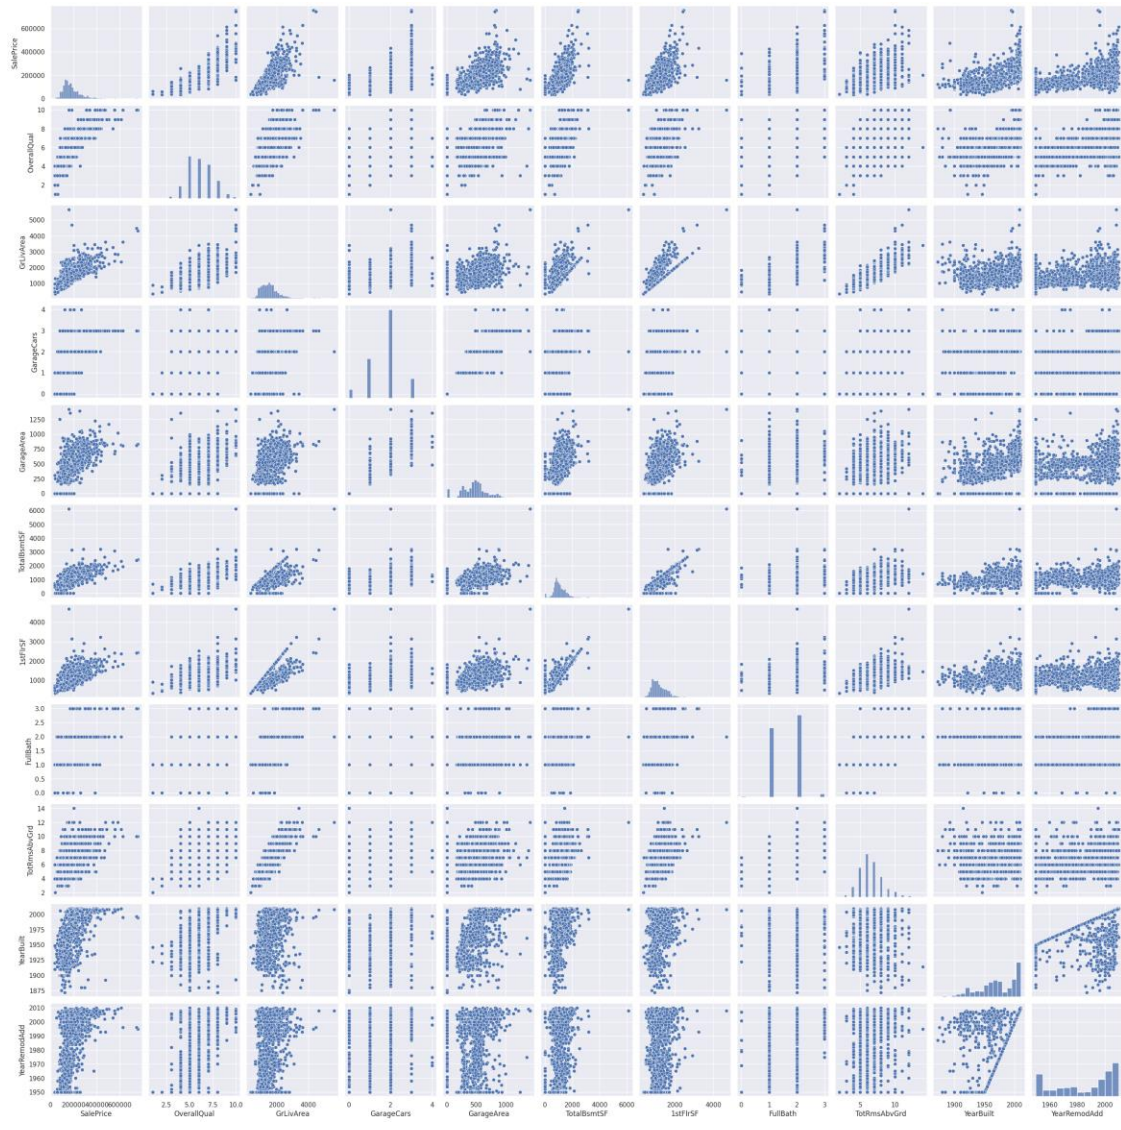
```

Scatter plot for numerical features that have high correlation with target.

```

sns.set()
numerical_features = ['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars', 'GarageArea', 'TotalBs
mtSF', '1stFlrSF', 'FullBath', 'TotRmsAbvGrd', 'YearBuilt', 'YearRemodAdd']
sns.pairplot(numerical_house_df[numerical_features], height= 2.5)
plt.show()

```


###Data Pre-Processing

#####Handling Missing Values "Imputation"

```
# Calculate missing data
missing_values = house_df.isnull().sum().sort_values(ascending=False)
missing_percent = (house_df.isnull().sum()/house_df.isnull().count()).sort_values(ascending=False)
missing_data = pd.concat([missing_values, missing_percent], axis=1, keys=['missing_values', 'missing_percent'])
missing_data.head(30)
```

	missing_values	missing_percent
PoolQC	1453	0.995205
MiscFeature	1406	0.963014
Alley	1369	0.937671
Fence	1179	0.807534
FireplaceQu	690	0.472603
LotFrontage	259	0.177397
GarageYrBlt	81	0.055479
GarageCond	81	0.055479
GarageType	81	0.055479
GarageFinish	81	0.055479
GarageQual	81	0.055479
BsmtExposure	38	0.026027
BsmtFinType2	38	0.026027
BsmtCond	37	0.025342
BsmtQual	37	0.025342
BsmtFinType1	37	0.025342
MasVnrArea	8	0.005479
MasVnrType	8	0.005479
Electrical	1	0.000685
MSSubClass	0	0.000000
Fireplaces	0	0.000000
Functional	0	0.000000
KitchenQual	0	0.000000
KitchenAbvGr	0	0.000000
BedroomAbvGr	0	0.000000
HalfBath	0	0.000000
FullBath	0	0.000000
BsmtHalfBath	0	0.000000
TotRmsAbvGrd	0	0.000000
GarageCars	0	0.000000

```
# bar chart for missing values
missing_values = missing_values[missing_values > 0]
missing_values.plot.bar()
plt.show()
```



data type of missing values in each feature.

```
house_df.loc[:, missing_values.index].dtypes
```

```
PoolQC          object
MiscFeature     object
Alley           object
Fence           object
FireplaceQu     object
LotFrontage     float64
GarageYrBlt     float64
GarageCond      object
GarageType      object
GarageFinish    object
GarageQual      object
BsmtExposure    object
BsmtFinType2    object
BsmtCond        object
BsmtQual        object
BsmtFinType1    object
MasVnrArea      float64
MasVnrType      object
Electrical      object
dtype: object
```

```

# Replace missing values in Fire place quality with None.
house_df['FireplaceQu'].fillna('None',inplace = True)# No fireplace
house_df['GarageCond'].fillna('None',inplace = True)# No Garage
house_df['GarageType'].fillna('None',inplace = True) # No Garage
house_df['GarageFinish'].fillna('None',inplace = True)# No Garage
house_df['GarageQual'].fillna('None',inplace = True)# No Garage
house_df['PoolQC'].fillna('None',inplace = True) # No Pool
house_df['MiscFeature'].fillna('None',inplace = True)# No fireplace
house_df['Alley'].fillna('None',inplace = True)# No Alley
house_df['Fence'].fillna('None',inplace = True)# No Fence
house_df['BsmtExposure'].fillna('None',inplace = True)# No Basement
house_df['BsmtFinType2'].fillna('None',inplace = True)# No Basement
house_df['BsmtCond'].fillna('None',inplace = True)# No Basement
house_df['BsmtQual'].fillna('None',inplace = True)# No Basement
house_df['BsmtFinType1'].fillna('None',inplace = True)# No Basement

# Replace missing values with mode.
house_df['MasVnrType']=house_df['MasVnrType'].fillna(house_df['MasVnrType'].mode()[0])
house_df['Electrical']=house_df['Electrical'].fillna(house_df['Electrical'].mode()[0])

# Replace missing values from numerical columns.
house_df['LotFrontage']=house_df['LotFrontage'].fillna(house_df['LotFrontage'].median())
house_df['GarageYrBlt']= house_df['GarageYrBlt'].fillna(house_df['GarageYrBlt'].median())
house_df['MasVnrArea'] = house_df['MasVnrArea'].fillna(house_df['MasVnrArea'].median())

house_df.isnull().sum().sort_values(ascending = False).head(20)

MSSubClass      0
MSZoning         0
GarageYrBlt      0
GarageType       0
FireplaceQu      0
Fireplaces       0
Functional       0
TotRmsAbvGrd     0
KitchenQual      0
KitchenAbvGr     0
BedroomAbvGr     0
HalfBath         0
FullBath         0
BsmtHalfBath     0
BsmtFullBath     0
GrLivArea        0
LowQualFinSF     0
2ndFlrSF         0
1stFlrSF         0
GarageFinish     0
dtype: int64

```


#####Log transofmation

```

# Sale Price before Log transformation
sns.histplot(house_df['SalePrice'], kde = True)
plt.title("Distribution of SalePrice")

Text(0.5, 1.0, 'Distribution of SalePrice')

```



```
# Target Log transformation
house_df['SalePrice'] = np.log(house_df['SalePrice'])

# SalePrice variance after Log transformation
house_df['SalePrice'].var()

0.15956179505733453

# SalePrice after transformation
sns.histplot(house_df['SalePrice'], kde = True)
plt.title("Distribution of SalePrice")

Text(0.5, 1.0, 'Distribution of SalePrice')
```



Handling Outliers

#numerica columns

numerical_house_df.columns

```
Index(['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
       'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2',
       'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
       'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
       'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
       'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
       'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',
       'MoSold', 'YrSold', 'SalePrice'],
      dtype='object')
```

boxplot for numerical features before deleting outliers

num_features = len(numerical_house_df.columns)

Calculate the figure size

fig, axes = plt.subplots(nrows= 8, ncols= 5, figsize=(15, 40))

Flatten the axes array to simplify indexing

axes = axes.flatten()

Iterate over numerical features and create individual box plots in subplots

```
for i, column in enumerate(numerical_house_df.columns):
```

```
    if i < num_features:
```

```
        sns.boxplot(x=house_df[column], ax=axes[i])
```

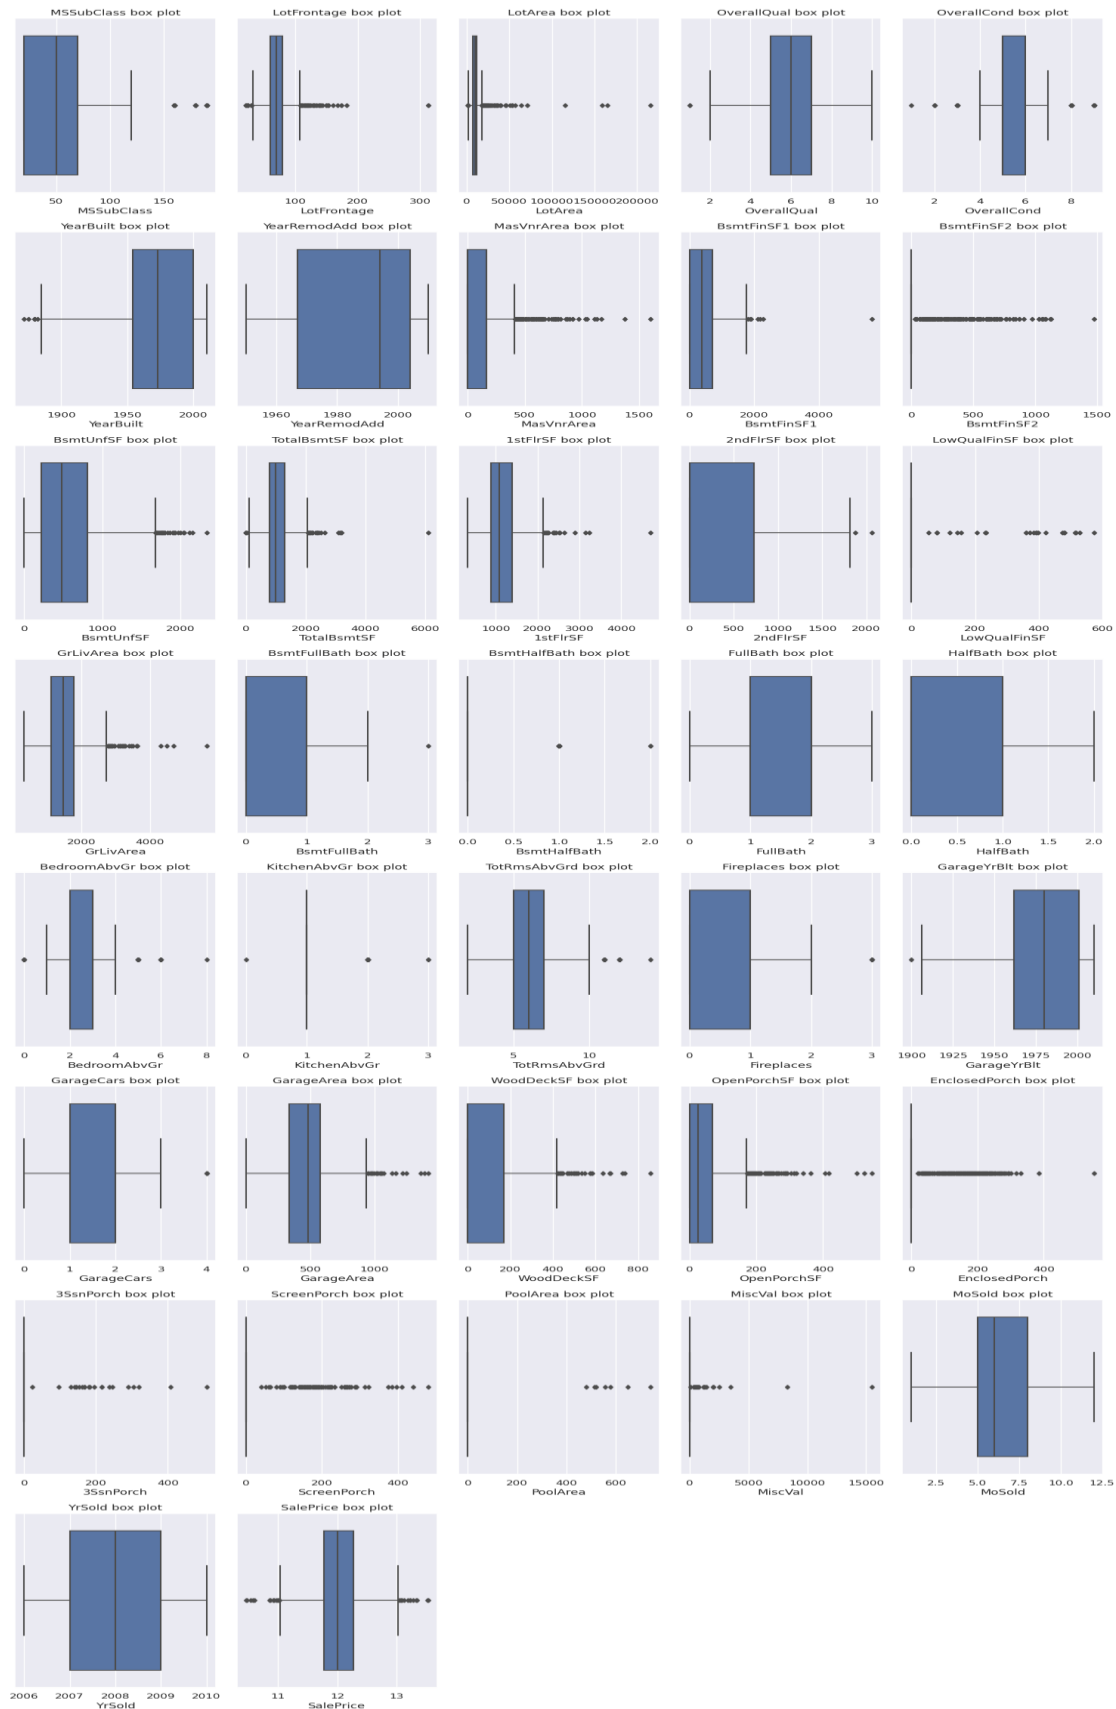
```
        axes[i].set_title(f'{column} box plot')
```

Minimize to only 37 plot

```
for i in range(len(numerical_house_df.columns), len(axes)):
```

```
    fig.delaxes(axes[i])
```

```
#prevent overlapping title.  
plt.tight_layout()  
  
plt.show()
```




```

# Create function to calculate upper and lower whisker for each numerical feature.
def capping(df, cols, factor):
    for col in cols:
        Q1 = df[col].quantile(0.25) # Calculate first quartile
        Q3 = df[col].quantile(0.75) # calculate second quartile
        IQR = Q3 - Q1 # Interquartile
        upper = Q3 + (factor*IQR) # upper_whisker
        lower = Q1 - (factor*IQR) # lower_whisker
        df[col] = np.where(df[col]>upper, upper, np.where(df[col]<lower, lower, df[col])) # change outliers

# Handle outliers with capping method.
capping(house_df, house_df.select_dtypes(include=['int64', 'float64']), 1.5)

# boxplot after deleting outliers
num_features = len(numerical_house_df.columns)
# figure size
fig, axes = plt.subplots(nrows= 8, ncols= 5, figsize=(15, 40))
# Flatten axes to simplify indexing
axes = axes.flatten()
# create individual box plots for each feature in subplots
for i, column in enumerate(numerical_house_df.columns):
    if i < num_features:
        sns.boxplot(x=house_df[column], ax=axes[i])
        axes[i].set_title(f'{column} box plot')
#minimize to only 37 plot
for i in range(len(numerical_house_df.columns), len(axes)):
    fig.delaxes(axes[i])
#prevent overlapping title.
plt.tight_layout()
plt.show()

```



####Handling Categorical Features

Categorical columns

categorical_house_df.columns

```
Index(['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities',  
      'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',  
      'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMat1', 'Exterior1st',  
      'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation',  
      'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',  
      'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual',  
      'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual',  
      'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature',  
      'SaleType', 'SaleCondition'],  
      dtype='object')
```

Numerical columns

numerical_house_df.columns

```
Index(['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',  
      'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2',  
      'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',  
      'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',  
      'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',  
      'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',  
      'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',  
      'MoSold', 'YrSold', 'SalePrice'],  
      dtype='object')
```

Categorical encoding

label = LabelEncoder()

for col in house_df.columns:

if house_df[col].dtype == 'object':

house_df[col] = label.fit_transform(house_df[col])

house_df.shape

(1460, 80)

house_df.head()

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	60.0	3	65.0	8450.0	1	1	3	
1	20.0	3	80.0	9600.0	1	1	3	
2	60.0	3	68.0	11250.0	1	1	0	
3	70.0	3	60.0	9550.0	1	1	0	
4	60.0	3	84.0	14260.0	1	1	0	

	LandContour	Utilities	LotConfig	...	PoolArea	PoolQC	Fence	\
0	3	0	4	...	0.0	3	4	
1	3	0	2	...	0.0	3	4	
2	3	0	4	...	0.0	3	4	
3	3	0	0	...	0.0	3	4	
4	3	0	2	...	0.0	3	4	

	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	1	0.0	2.0	2008.0	8	4	12.247694
1	1	0.0	5.0	2007.0	8	4	12.109011
2	1	0.0	9.0	2008.0	8	4	12.317167
3	1	0.0	2.0	2006.0	8	0	11.849398
4	1	0.0	12.0	2008.0	8	4	12.429216

```
[5 rows x 80 columns]
```

###Splitting to train and test

```
X = house_df.drop('SalePrice', axis = 1)
y = house_df['SalePrice']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
X_train
```

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
254	20.0	3	70.0	8400.0	1	1	3	
1066	60.0	3	59.0	7837.0	1	1	0	
638	30.0	3	67.0	8777.0	1	1	3	
799	50.0	3	60.0	7200.0	1	1	3	
380	50.0	3	50.0	5000.0	1	2	3	
...	
1095	20.0	3	78.0	9317.0	1	1	0	
1130	50.0	3	65.0	7804.0	1	1	3	
1294	20.0	3	60.0	8172.0	1	1	3	
860	50.0	3	55.0	7642.0	1	1	3	
1126	120.0	3	53.0	3684.0	1	1	3	

	LandContour	Utilities	LotConfig	...	ScreenPorch	PoolArea	PoolQC	\
254	3	0	4	...	0.0	0.0	3	
1066	3	0	4	...	0.0	0.0	3	
638	3	0	4	...	0.0	0.0	3	
799	3	0	0	...	0.0	0.0	3	
380	3	0	4	...	0.0	0.0	3	
...	
1095	3	0	4	...	0.0	0.0	3	
1130	3	0	4	...	0.0	0.0	3	
1294	3	0	4	...	0.0	0.0	3	
860	3	0	0	...	0.0	0.0	3	
1126	3	0	4	...	0.0	0.0	3	

	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition
254	4	1	0.0	6.0	2010.0	8	4
1066	4	1	0.0	5.0	2009.0	8	4
638	2	1	0.0	5.0	2008.0	8	4
799	2	1	0.0	6.0	2007.0	8	4
380	4	1	0.0	5.0	2010.0	8	4
...
1095	4	1	0.0	3.0	2007.0	8	4
1130	2	1	0.0	12.0	2009.0	8	4
1294	4	1	0.0	4.0	2006.0	8	4
860	0	1	0.0	6.0	2007.0	8	4
1126	4	1	0.0	6.0	2009.0	8	4

```
[1168 rows x 79 columns]
```

#####Scaling the Data

```
scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
X_train_scaled = pd.DataFrame(X_train_scaled, columns=X_train.columns)
X_train_scaled
```

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	\
0	-0.933485	-0.054796	0.044442	-0.369609	0.058621	0.047674	
1	0.135447	-0.054796	-0.592466	-0.527046	0.058621	0.047674	
2	-0.666252	-0.054796	-0.129260	-0.264185	0.058621	0.047674	
3	-0.131786	-0.054796	-0.534566	-0.705176	0.058621	0.047674	
4	-0.131786	-0.054796	-1.113573	-1.320382	0.058621	4.025068	
...	
1163	-0.933485	-0.054796	0.507648	-0.113180	0.058621	0.047674	
1164	-0.131786	-0.054796	-0.245062	-0.536274	0.058621	0.047674	
1165	-0.933485	-0.054796	-0.534566	-0.433367	0.058621	0.047674	
1166	-0.131786	-0.054796	-0.824069	-0.581576	0.058621	0.047674	
1167	1.738845	-0.054796	-0.939871	-1.688388	0.058621	0.047674	

	LotShape	LandContour	Utilities	LotConfig	...	ScreenPorch	PoolArea	\
0	0.765535	0.299798	-0.029273	0.630128	...	0.0	0.0	
1	-1.356644	0.299798	-0.029273	0.630128	...	0.0	0.0	
2	0.765535	0.299798	-0.029273	0.630128	...	0.0	0.0	
3	0.765535	0.299798	-0.029273	-1.792884	...	0.0	0.0	
4	0.765535	0.299798	-0.029273	0.630128	...	0.0	0.0	
...	
1163	-1.356644	0.299798	-0.029273	0.630128	...	0.0	0.0	
1164	0.765535	0.299798	-0.029273	0.630128	...	0.0	0.0	
1165	0.765535	0.299798	-0.029273	0.630128	...	0.0	0.0	
1166	0.765535	0.299798	-0.029273	-1.792884	...	0.0	0.0	
1167	0.765535	0.299798	-0.029273	0.630128	...	0.0	0.0	

	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType	\
0	0.066503	0.467615	-0.190299	0.0	-0.133417	1.650065	0.316662	
1	0.066503	0.467615	-0.190299	0.0	-0.508010	0.893677	0.316662	
2	0.066503	-1.343907	-0.190299	0.0	-0.508010	0.137290	0.316662	
3	0.066503	-1.343907	-0.190299	0.0	-0.133417	-0.619098	0.316662	
4	0.066503	0.467615	-0.190299	0.0	-0.508010	1.650065	0.316662	
...	
1163	0.066503	0.467615	-0.190299	0.0	-1.257196	-0.619098	0.316662	
1164	0.066503	-1.343907	-0.190299	0.0	2.114141	0.893677	0.316662	
1165	0.066503	0.467615	-0.190299	0.0	-0.882603	-1.375486	0.316662	
1166	0.066503	-3.155429	-0.190299	0.0	-0.133417	-0.619098	0.316662	
1167	0.066503	0.467615	-0.190299	0.0	-0.133417	0.893677	0.316662	

	SaleCondition
0	0.201772
1	0.201772
2	0.201772
3	0.201772
4	0.201772
...	...
1163	0.201772
1164	0.201772
1165	0.201772
1166	0.201772
1167	0.201772

```
[1168 rows x 79 columns]
```

```
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X_test.columns)
X_test_scaled
```

```

MSSubClass MSZoning LotFrontage LotArea Street Alley \
0 -0.933485 -0.054796 0.044442 -0.365694 0.058621 0.047674
1 0.135447 -0.054796 1.665663 0.708680 0.058621 0.047674
2 -0.666252 1.586277 -0.766169 -0.213011 0.058621 -3.929719
3 -0.131786 1.586277 -1.113573 -1.320382 0.058621 0.047674
4 -0.933485 -0.054796 1.144556 0.888208 0.058621 0.047674
.. ...
287 -0.666252 1.586277 -1.113573 -1.061716 0.058621 0.047674
288 -0.933485 -0.054796 2.215720 1.799832 0.058621 0.047674
289 0.135447 -0.054796 -0.360863 -0.425817 0.058621 0.047674
290 0.402680 -0.054796 -0.534566 -0.178336 0.058621 0.047674
291 -0.933485 -0.054796 0.044442 -0.447908 0.058621 0.047674

LotShape LandContour Utilities LotConfig ... ScreenPorch PoolArea \
0 0.765535 0.299798 -0.029273 0.630128 ... 0.0 0.0
1 -1.356644 0.299798 -0.029273 -1.792884 ... 0.0 0.0
2 0.765535 0.299798 -0.029273 0.630128 ... 0.0 0.0
3 0.765535 0.299798 -0.029273 -1.792884 ... 0.0 0.0
4 -1.356644 -2.618233 -0.029273 0.630128 ... 0.0 0.0
.. ...
287 0.765535 -4.077249 -0.029273 0.630128 ... 0.0 0.0
288 -1.356644 -1.159218 -0.029273 0.630128 ... 0.0 0.0
289 0.765535 0.299798 -0.029273 0.630128 ... 0.0 0.0
290 0.765535 0.299798 -0.029273 0.630128 ... 0.0 0.0
291 0.765535 0.299798 -0.029273 0.630128 ... 0.0 0.0

PoolQC Fence MiscFeature MiscVal MoSold YrSold SaleType \
0 0.066503 -1.343907 -0.190299 0.0 -1.631789 -1.375486 0.316662
1 0.066503 0.467615 -0.190299 0.0 -0.882603 1.650065 0.316662
2 0.066503 0.467615 -0.190299 0.0 -1.257196 1.650065 0.316662
3 0.066503 0.467615 -0.190299 0.0 1.364955 -1.375486 0.316662
4 0.066503 0.467615 -0.190299 0.0 0.990362 0.893677 0.316662
.. ...
287 0.066503 -1.343907 -0.190299 0.0 -1.257196 -0.619098 0.316662
288 0.066503 0.467615 -0.190299 0.0 -0.133417 0.893677 0.316662
289 0.066503 0.467615 -0.190299 0.0 1.364955 0.137290 0.316662
290 0.066503 -1.343907 -0.190299 0.0 1.364955 0.893677 0.316662
291 0.066503 0.467615 -0.190299 0.0 0.241176 0.893677 0.316662

SaleCondition
0 0.201772
1 0.201772
2 0.201772
3 0.201772
4 0.201772
.. ...
287 -1.653892
288 0.201772
289 0.201772
290 0.201772
291 0.201772

```

[292 rows x 79 columns]

#####Modeling and Hyperparameter Tuning

#####LASSO

```

# create model instance
lasso_model = Lasso(alpha = 0.001, random_state = 10) #create model instance

```

```

# Use RandomizedSearchCV
lasso_parameters = {'alpha': [0.001, 0.01, 0.1, 0.000391, 0.000491, 1, 10, 100, 1000],
                    'random_state': [1, 5, 10, 24, 42]}
lasso_model_random = RandomizedSearchCV(lasso_model, lasso_parameters, n_iter=10, cv=5, scoring
= 'neg_mean_absolute_error', n_jobs = -1)
lasso_model_random.fit(X_train_scaled, y_train)

RandomizedSearchCV(cv=5, estimator=Lasso(alpha=0.001, random_state=10),
                    n_jobs=-1,
                    param_distributions={'alpha': [0.001, 0.01, 0.1, 0.000391,
                                                    0.000491, 1, 10, 100, 1000],
                                          'random_state': [1, 5, 10, 24, 42]},
                    scoring='neg_mean_absolute_error')

# Lasso validation
lasso_r2_cv = cross_val_score(lasso_model, X_train_scaled, y_train, scoring="r2", cv = 5)
lasso_mae_cv = -cross_val_score(lasso_model, X_train_scaled, y_train, scoring="neg_mean_absolute
_error", cv = 5)
lasso_rmse_cv = -cross_val_score(lasso_model, X_train_scaled, y_train, scoring="neg_root_mean_sq
uared_error", cv = 5)
# Print score and scoring means
print("R2 scores:", lasso_r2_cv)
print("Mean R2:", np.mean(lasso_r2_cv))
print("MAE scores:", lasso_mae_cv)
print("Mean MAE:", np.mean(lasso_mae_cv))
print("RSME scores:", lasso_rmse_cv)
print("Mean RSME:", np.mean(lasso_rmse_cv))

R2 scores: [0.91514172 0.86114004 0.88331953 0.90540684 0.90613434]
Mean R2: 0.8942284939023724
MAE scores: [0.0812163 0.0917392 0.09074863 0.08588829 0.07771177]
Mean MAE: 0.0854608399317327
RSME scores: [0.10970036 0.13920932 0.13993263 0.11722237 0.1083503 ]
Mean RSME: 0.12288299421975364

# figure out best hyperparameters
print(lasso_model_random.best_estimator_)

Lasso(alpha=0.000391, random_state=10)

# Training LinearRegression model on training set.
lasso_model.fit(X_train_scaled, y_train)

Lasso(alpha=0.001, random_state=10)

# Model Prediction
lasso_y_pred = lasso_model.predict(X_test_scaled)

# Calculate R2, MAE, RMSE
lasso_r2_score = r2_score(y_test, lasso_y_pred)
lasso_mae_score = mean_absolute_error(y_test, lasso_y_pred)
lasso_rmse_score = np.sqrt(mean_squared_error(y_test, lasso_y_pred))
# Print scores
print("R2 score:", lasso_r2_score)
print("MAE scores:", lasso_mae_score)
print("RMSE score:", lasso_rmse_score)

R2 score: 0.9044818230200362
MAE scores: 0.08970264159372936
RMSE score: 0.12698196744804813

```

#####KNN Regressor

```
# KNN hyperparameter tuning.
knn_model = KNeighborsRegressor()
knn_parameters = {
    'n_neighbors': np.arange(5,20)
}
knn_model_random = RandomizedSearchCV(knn_model, knn_parameters, n_iter=10, cv=5, random_state
=42, scoring = 'neg_mean_absolute_error', n_jobs = -1)
knn_model_random.fit(X_train_scaled, y_train)

RandomizedSearchCV(cv=5, estimator=KNeighborsRegressor(), n_jobs=-1,
    param_distributions={'n_neighbors': array([ 5,  6,  7,  8,  9, 10, 11, 12,
13, 14, 15, 16, 17, 18, 19])},
    random_state=42, scoring='neg_mean_absolute_error')

# figure out best hyperparameters
print(knn_model_random.best_estimator_)

KNeighborsRegressor(n_neighbors=7)

# KNN Cross-Validation
knn_model = KNeighborsRegressor(7)
knn_r2_cv = cross_val_score(knn_model, X_train_scaled, y_train, scoring="r2", cv = 5)
knn_mae_cv = -cross_val_score(knn_model, X_train_scaled, y_train, scoring="neg_mean_absolute_err
or", cv = 5)
knn_rmse_cv = -cross_val_score(knn_model, X_train_scaled, y_train, scoring="neg_root_mean_square
d_error", cv = 5)
# Print score and scoring means
print("R2 scores:", knn_r2_cv)
print("Mean R2:", np.mean(knn_r2_cv))
print("MAE scores:", knn_mae_cv)
print("Mean MAE:", np.mean(knn_mae_cv))
print("RSME scores:", knn_rmse_cv)
print("Mean RSME:", np.mean(knn_rmse_cv))

R2 scores: [0.82942321 0.80857873 0.78376055 0.80545169 0.81547797]
Mean R2: 0.8085384293636986
MAE scores: [0.11217005 0.11960553 0.13581658 0.1229724  0.11066458]
Mean MAE: 0.12024582799950663
RSME scores: [0.1555324  0.16344623 0.19049666 0.16811032 0.15191507]
Mean RSME: 0.1659001348252332

# KNN fitting to data.
knn_model.fit(X_train_scaled, y_train)

KNeighborsRegressor(n_neighbors=7)

# KNN prediction
knn_y_pred = knn_model.predict(X_test_scaled)

# Calculate and print scores
knn_r2_score = r2_score(y_test, knn_y_pred)
knn_mae_score = mean_absolute_error(y_test, knn_y_pred)
knn_rmse_score = np.sqrt(mean_squared_error(y_test, knn_y_pred))
# Print scores
print("R2 score:", knn_r2_score)
print("MAE scores:", knn_mae_score)
print("RMSE score:", knn_rmse_score)
```



```
R2 score: 0.8188392441996817
MAE scores: 0.1286051517763571
RMSE score: 0.17487635843662072
```

#####SVR

```
# SVR cross_validation
svr_model = SVR()# create model instance
svr_r2_cv = cross_val_score(svr_model,X_train_scaled,y_train, scoring="r2",cv = 5)
svr_mae_cv = -cross_val_score(svr_model,X_train_scaled,y_train, scoring="neg_mean_absolute_err
or",cv = 5)
svr_rmse_cv = -cross_val_score(svr_model,X_train_scaled,y_train, scoring="neg_root_mean_squar
e_d_error",cv = 5)
# Print score and scoring means
print("R2 scores:", svr_r2_cv)
print("Mean R2:", np.mean(svr_r2_cv))
print("MAE scores:", svr_mae_cv)
print("Mean MAE:", np.mean(svr_mae_cv))
print("RSME scores:", svr_rmse_cv)
print("Mean RSME:", np.mean(svr_rmse_cv))
```

```
R2 scores: [0.85125548 0.85978534 0.84081795 0.82259625 0.84384907]
Mean R2: 0.8436608180793096
MAE scores: [0.09320769 0.10188274 0.11181347 0.10880024 0.09787004]
Mean MAE: 0.10271483449661556
RSME scores: [0.14523838 0.13988673 0.16344323 0.16053216 0.1397491 ]
Mean RSME: 0.1497699199056675
```

```
# fit SVR model
svr_model.fit(X_train_scaled, y_train)
```

```
SVR()
```

```
# Make predictions on the test data.
svr_y_pred = svr_model.predict(X_test_scaled)
```

```
# Calculate and print evaluations scores
svr_r2_score = r2_score(y_test, svr_y_pred)
svr_mae_score = mean_absolute_error(y_test, svr_y_pred)
svr_rmse_score = np.sqrt(mean_squared_error(y_test, svr_y_pred))
# Print the evaluation metrics.
print("R2 score:", svr_r2_score)
print("MAE scores:", svr_mae_score)
print("RMSE score:", svr_rmse_score)
```

```
R2 score: 0.821242610879984
MAE scores: 0.1110363341623632
RMSE score: 0.17371248802626518
```

#####Decision Tree Regressor

```
# DTR hyperparameter tuning using RandomizedSearch cv
dtr_model = DecisionTreeRegressor() #create model instance
```

```
dtr_parameters = {
    'max_depth': np.arange(1, 21),
    'min_samples_split': np.arange(2, 11),
    'min_samples_leaf': np.arange(1, 11),
    'max_features': [20,30,50]
}
dtr_model_random = RandomizedSearchCV(dtr_model, dtr_parameters, n_iter=10, cv=5, random_state
```

```

=42,scoring = 'neg_mean_absolute_error',n_jobs = -1)
dtr_model_random.fit(X_train_scaled,y_train)

RandomizedSearchCV(cv=5, estimator=DecisionTreeRegressor(), n_jobs=-1,
                    param_distributions={'max_depth': array([ 1,  2,  3,  4,  5,  6,  7,  8,  9
, 10, 11, 12, 13, 14, 15, 16, 17,
                    18, 19, 20]),
                    'max_features': [20, 30, 50],
                    'min_samples_leaf': array([ 1,  2,  3,  4,  5,  6,  7,
8,  9, 10]),
                    'min_samples_split': array([ 2,  3,  4,  5,  6,  7,  8
,  9, 10])},
                    random_state=42, scoring='neg_mean_absolute_error')

# figure out best hyperparameters
print(dtr_model_random.best_estimator_)

DecisionTreeRegressor(max_depth=20, max_features=50, min_samples_leaf=9,
                      min_samples_split=10)

# DTR validation
dtr_model = DecisionTreeRegressor(max_depth=11, min_samples_leaf=7, min_samples_split=4)
dtr_r2_cv = cross_val_score(dtr_model,X_train_scaled,y_train, scoring="r2",cv = 5)
dtr_mae_cv = -cross_val_score(dtr_model,X_train_scaled,y_train, scoring="neg_mean_absolute_err
or",cv = 5)
dtr_rmse_cv = -cross_val_score(dtr_model,X_train_scaled,y_train, scoring="neg_root_mean_square
d_error",cv = 5)
# Print score and scoring means
print("R2 scores:", dtr_r2_cv)
print("Mean R2:", np.mean(dtr_r2_cv))
print("MAE scores:", dtr_mae_cv)
print("Mean MAE:", np.mean(dtr_mae_cv))
print("RSME scores:", dtr_rmse_cv)
print("Mean RSME:", np.mean(dtr_rmse_cv))

R2 scores: [0.79393244 0.73275859 0.81215354 0.785475    0.76992605]
Mean R2: 0.778849125071273
MAE scores: [0.1257418  0.13655795 0.13387828 0.13613404 0.12706427]
Mean MAE: 0.13187526633460778
RSME scores: [0.17219127 0.19365437 0.1775503  0.17646496 0.16934539]
Mean RSME: 0.1778412568497989

# fitting dtr model.
dtr_model.fit(X_train_scaled,y_train)

DecisionTreeRegressor(max_depth=11, min_samples_leaf=7, min_samples_split=4)

# dtr prediction
dtr_y_pred = dtr_model.predict(X_test_scaled)

# calculate and print dtr scores
dtr_r2_score = r2_score(y_test, dtr_y_pred)
dtr_mae_score = mean_absolute_error(y_test, dtr_y_pred)
dtr_rmse_score = np.sqrt(mean_squared_error(y_test,dtr_y_pred))
# Print scores
print("R2 score:", dtr_r2_score)
print("MAE scores:", dtr_mae_score)
print("RMSE score:", dtr_rmse_score)

R2 score: 0.7911529039464633
MAE scores: 0.14403474799395483
RMSE score: 0.18776440251810755

```

#####Random Forest Regressor

```
# Rf hyperparameter tuning.
rf_model = RandomForestRegressor()
rf_parameters = {
    'min_samples_leaf':[10,15,20],
    'min_samples_split':[10,15,20],
    'max_depth': [4,6,8,10],
}
rf_model_random = RandomizedSearchCV(rf_model, rf_parameters, n_iter=10, cv=5, random_state=42,
scoring = 'neg_mean_absolute_error',n_jobs = -1)
rf_model_random.fit(X_train_scaled,y_train)

RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(), n_jobs=-1,
                    param_distributions={'max_depth': [4, 6, 8, 10],
                                         'min_samples_leaf': [10, 15, 20],
                                         'min_samples_split': [10, 15, 20]},
                    random_state=42, scoring='neg_mean_absolute_error')

# figure out best hyperparameters
print(rf_model_random.best_estimator_)

RandomForestRegressor(max_depth=8, min_samples_leaf=15, min_samples_split=10)

# cross validation to random forest
rf_model = RandomForestRegressor()
random_forest_r2_cv = cross_val_score(rf_model,X_train_scaled,y_train, scoring="r2",cv = 5)
random_forest_mae_cv = -cross_val_score(rf_model,X_train_scaled,y_train, scoring="neg_mean_absolute_error",cv = 5)
random_forest_rmse_cv = -cross_val_score(rf_model,X_train_scaled,y_train, scoring="neg_root_mean_squared_error",cv = 5)
# Print score and scoring means
print("R2 scores:", random_forest_r2_cv)
print("Mean R2:", np.mean(random_forest_r2_cv))
print("MAE scores:", random_forest_mae_cv)
print("Mean MAE:", np.mean(random_forest_mae_cv))
print("RSME scores:", random_forest_rmse_cv)
print("Mean RSME:", np.mean(random_forest_rmse_cv))

R2 scores: [0.89803974 0.82701911 0.87129568 0.86519458 0.88786007]
Mean R2: 0.8698818338757828
MAE scores: [0.08374276 0.10587958 0.10187934 0.10281325 0.08199719]
Mean MAE: 0.09526242445464875
RSME scores: [0.12076353 0.15440703 0.14723576 0.13856128 0.11617691]
Mean RSME: 0.1354289046572923

#Random Forest fit
rf_model.fit(X_train_scaled,y_train)

RandomForestRegressor()

# random forest prediction
rf_y_pred = rf_model.predict(X_test_scaled)

# calculate and print scores
rf_r2_score = r2_score(y_test, rf_y_pred)
rf_mae_score = mean_absolute_error(y_test, rf_y_pred)
rf_rmse_score = np.sqrt(mean_squared_error(y_test, rf_y_pred))
# Print scores
print("R2 score:", rf_r2_score)
```

```
print("MAE scores:", rf_mae_score)
print("RMSE score:", rf_rmse_score)
```

```
R2 score: 0.8899633870185536
MAE scores: 0.09728684639777124
RMSE score: 0.13629114689424793
```

#####Gradient Booster Regressor

Gradient Boosting Hyperparameter tuning

```
gbr_model = GradientBoostingRegressor()
gbr_parameters = {
    'n_estimators':[2,3,4,5,6,7,20,30,50],
    'learning_rate':[0.01, 0.03, 0.05, 0.1,0.001],
    'max_depth':[3, 4, 5, 6],
    'alpha':[0.001,0.1,0.5,0.2]
}
gbr_model_random = RandomizedSearchCV(gbr_model,gbr_parameters, n_iter=10, cv=5,scoring = 'neg
_mean_absolute_error',n_jobs = -1,error_score='raise')
gbr_model_random.fit(X_train_scaled,y_train)
```

```
RandomizedSearchCV(cv=5, error_score='raise',
                    estimator=GradientBoostingRegressor(), n_jobs=-1,
                    param_distributions={'alpha': [0.001, 0.1, 0.5, 0.2],
                                         'learning_rate': [0.01, 0.03, 0.05, 0.1,
                                                           0.001],
                                         'max_depth': [3, 4, 5, 6],
                                         'n_estimators': [2, 3, 4, 5, 6, 7, 20,
                                                         30, 50]},
                    scoring='neg_mean_absolute_error')
```

figure out best hyperparameters

```
print(gbr_model_random.best_estimator_)
```

```
GradientBoostingRegressor(alpha=0.5, learning_rate=0.05, max_depth=5,
                           n_estimators=30)
```

Gradient Booster cross validation

```
gbr_model = GradientBoostingRegressor(alpha=0.5, n_estimators=30) # create model instance
gbr_r2_cv = cross_val_score(gbr_model,X_train_scaled,y_train, scoring="r2",cv = 5)
gbr_mae_cv = -cross_val_score(gbr_model,X_train_scaled,y_train, scoring="neg_mean_absolute_err
or",cv = 5)
gbr_rmse_cv = -cross_val_score(gbr_model,X_train_scaled,y_train, scoring="neg_root_mean_squar
e_d_error",cv = 5)
# Print score and scoring means
print("R2 scores:", gbr_r2_cv)
print("Mean R2:", np.mean(gbr_r2_cv))
print("MAE scores:", gbr_mae_cv)
print("Mean MAE:", np.mean(gbr_mae_cv))
print("RSME scores:", gbr_rmse_cv)
print("Mean RSME:", np.mean(gbr_rmse_cv))
```

```
R2 scores: [0.8673904  0.83720537 0.83951086 0.8478362  0.87889375]
Mean R2: 0.8541673164693911
MAE scores: [0.09652216 0.11075329 0.11661752 0.10888996 0.08685095]
Mean MAE: 0.10392677621492344
RSME scores: [0.13704732 0.1519922  0.1641129  0.14898419 0.12283623]
Mean RSME: 0.1449945684809948
```

Fit gbr model to data

```
gbr_model.fit(X_train_scaled, y_train)
```

```

GradientBoostingRegressor(alpha=0.5, n_estimators=30)

# gbr prediction
gbr_y_pred = gbr_model.predict(X_test_scaled)

# Calculate and print evaluation scores
gbr_r2_score = r2_score(y_test, gbr_y_pred)
gbr_mae_score = mean_absolute_error(y_test, gbr_y_pred)
gbr_rmse_score = np.sqrt(mean_squared_error(y_test, gbr_y_pred))
# Print the evaluation metrics.
print("R2 score:", gbr_r2_score)
print("MAE scores:", gbr_mae_score)
print("RMSE score:", gbr_rmse_score)

R2 score: 0.8776502352477635
MAE scores: 0.10467667222559117
RMSE score: 0.14371450669300415

```

#####XGBoost

```

# XGBoost hyperparameter tuning
xgb_model = XGBRegressor()
xgb_parameters = {
    'n_estimators': [100, 200, 300, 500, 2200, 220],
    'learning_rate': [0.01, 0.03, 0.05, 0.1, 0.001],
    'max_depth': [3, 4, 5, 6],
    'min_child_weight': [1, 2, 3, 4, 1.7817, 0.178, 0.0178],
    'gamma': [0, 0.1, 0.2, 0.3, 0.005, 0.0045, 0.35, 0.5],
    'subsample': [0.6, 0.7, 0.8, 0.9],
    'colsample_bytree': [0.6, 0.7, 0.8, 0.9],
    'alpha': [0, 0.1, 0.5, 1]
}
xgb_model_random = RandomizedSearchCV(xgb_model, xgb_parameters, n_iter=10, cv=5, scoring = 'neg
_mean_absolute_error', n_jobs = -1)
xgb_model_random.fit(X_train_scaled, y_train)

RandomizedSearchCV(cv=5,
                    estimator=XGBRegressor(base_score=None, booster=None,
                                           callbacks=None,
                                           colsample_bylevel=None,
                                           colsample_bynode=None,
                                           colsample_bytree=None, device=None,
                                           early_stopping_rounds=None,
                                           enable_categorical=False,
                                           eval_metric=None, feature_types=None,
                                           gamma=None, grow_policy=None,
                                           importance_type=None,
                                           interaction_constraints=None,
                                           learning_rate=...,
                                           n_jobs=-1,
                                           param_distributions={'alpha': [0, 0.1, 0.5, 1],
                                                                  'colsample_bytree': [0.6, 0.7, 0.8,
                                                                                       0.9],
                                                                  'gamma': [0, 0.1, 0.2, 0.3, 0.005,
                                                                                       0.0045, 0.35, 0.5],
                                                                  'learning_rate': [0.01, 0.03, 0.05, 0.1,
                                                                                       0.001],
                                                                  'max_depth': [3, 4, 5, 6],
                                                                  'min_child_weight': [1, 2, 3, 4, 1.7817,
                                                                                       0.178, 0.0178],
                                                                  'n_estimators': [100, 200, 300, 500,

```

```

                2200, 220],
                'subsample': [0.6, 0.7, 0.8, 0.9]},
                scoring='neg_mean_absolute_error')

# figure out best hyperparameters
print(xgb_model_random.best_estimator_)

XGBRegressor(alpha=0.1, base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.8, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=0.0045, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.03, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=6, max_leaves=None,
              min_child_weight=1.7817, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=200, n_jobs=None,
              num_parallel_tree=None, ...)

# XGBoost validation
xgb_model = XGBRegressor( gamma = 0,n_estimators=220,learning_rate=0.05,random_state=42,max_de
pth=5,colsample_bytree=0.7,min_child_weight=1.7817,alpha=0.5)# create model instance
xgb_r2_cv = cross_val_score(xgb_model,X_train_scaled,y_train, scoring="r2",cv = 5)
xgb_mae_cv = -cross_val_score(xgb_model,X_train_scaled,y_train, scoring="neg_mean_absolute_err
or",cv = 5)
xgb_rmse_cv = -cross_val_score(xgb_model,X_train_scaled,y_train, scoring="neg_root_mean_square
d_error",cv = 5)
# Print score and scoring means
print("R2 scores:", xgb_r2_cv)
print("Mean R2:", np.mean(xgb_r2_cv))
print("MAE scores:", xgb_mae_cv)
print("Mean MAE:", np.mean(xgb_mae_cv))
print("RSME scores:", xgb_rmse_cv)
print("Mean RSME:", np.mean(xgb_rmse_cv))

R2 scores: [0.90824463 0.86600988 0.89993456 0.8837685  0.90858287]
Mean R2: 0.8933080880133396
MAE scores: [0.07781259 0.09374167 0.09168381 0.09632367 0.07593342]
Mean MAE: 0.08709903241328876
RSME scores: [0.11407137 0.13674649 0.12958716 0.12993991 0.10692778]
Mean RSME: 0.12345453994917192

# XGBoost fit to data.
xgb_model.fit(X_train_scaled,y_train)

XGBRegressor(alpha=0.5, base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.7, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=0, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.05, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=5, max_leaves=None,
              min_child_weight=1.7817, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=220, n_jobs=None,
              num_parallel_tree=None, ...)

#XGBoost prediction
xgb_y_pred = xgb_model.predict(X_test_scaled)

```

```

# calculate and print scores.
xgb_r2_score = r2_score(y_test, xgb_y_pred)
xgb_mae_score = mean_absolute_error(y_test, xgb_y_pred)
xgb_rmse_score = np.sqrt(mean_squared_error(y_test, xgb_y_pred))
# Print scores
print("R2 score:", xgb_r2_score)
print("MAE scores:", xgb_mae_score)
print("RMSE score:", xgb_rmse_score)

```

```

R2 score: 0.9137570160601971
MAE scores: 0.08521061986653618
RMSE score: 0.12065933469866713

```

#####LightGBM

```

# LGB Hyperparameter
lgb_model = lgb.LGBMRegressor()
lgb_parameters = {
    'n_estimators': [6,7,10,11,50,60],
    'learning_rate': [0.7,0.8,0.1,0.001,0.2,0.0003,0.3,0.5],
    'num_leaves': [5,6,7],
    'max_depth': [3, 4, 5, 6,50],
    'reg_alpha': [0, 0.1, 0.5, 1],
}
lgb_model_random = RandomizedSearchCV(lgb_model,lgb_parameters,n_iter=10, cv=5,scoring = 'neg_mean_absolute_error',n_jobs = -1)
lgb_model_random.fit(X_train_scaled,y_train)

```

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000574 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2926

[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 65

[LightGBM] [Info] Start training from score 12.031035

```

RandomizedSearchCV(cv=5, estimator=LGBMRegressor(), n_jobs=-1,
                  param_distributions={'learning_rate': [0.7, 0.8, 0.1, 0.001,
                                                         0.2, 0.0003, 0.3,
                                                         0.5],
                                     'max_depth': [3, 4, 5, 6, 50],
                                     'n_estimators': [6, 7, 10, 11, 50, 60],
                                     'num_leaves': [5, 6, 7],
                                     'reg_alpha': [0, 0.1, 0.5, 1]},
                  scoring='neg_mean_absolute_error')

```

figure out best hyperparameters

```
print(lgb_model_random.best_estimator_)
```

```

LGBMRegressor(learning_rate=0.3, max_depth=6, n_estimators=50, num_leaves=7,
              reg_alpha=1)

```

LightGBM validation

```

lgb_model = lgb.LGBMRegressor(learning_rate=0.2, max_depth=6, n_estimators=50, num_leaves=6,reg_alpha=0)# create model instance
lgb_r2_cv = cross_val_score(lgb_model,X_train_scaled,y_train, scoring="r2",cv = 5)
lgb_mae_cv = -cross_val_score(lgb_model,X_train_scaled,y_train, scoring="neg_mean_absolute_error",cv = 5)
lgb_rmse_cv = -cross_val_score(lgb_model,X_train_scaled,y_train, scoring="neg_root_mean_squared_error",cv = 5)

```

```

print("R2 scores:", lgb_r2_cv)
print("Mean R2:", np.mean(lgb_r2_cv))
print("MAE scores:", lgb_mae_cv)
print("Mean MAE:", np.mean(lgb_mae_cv))
print("RSME scores:", lgb_rmse_cv)
print("Mean RSME:", np.mean(lgb_rmse_cv))

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000357
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2828
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.030302
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000405
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.020671
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000375
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2822
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 64
[LightGBM] [Info] Start training from score 12.041320
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000423
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.034734
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000386
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2835
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.028148
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000417
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2828
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.030302
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000577
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.020671
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000382
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2822

```



```

[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 64
[LightGBM] [Info] Start training from score 12.041320
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000416
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.034734
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000405
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2835
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.028148
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000383
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2828
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.030302
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000374
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.020671
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000399
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2822
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 64
[LightGBM] [Info] Start training from score 12.041320
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000379
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.034734
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000438
seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2835
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.028148
R2 scores: [0.90362198 0.86798016 0.88655032 0.88880053 0.89486506]
Mean R2: 0.8883636114462181
MAE scores: [0.08291353 0.08959073 0.09415874 0.09406981 0.07951395]
Mean MAE: 0.08804935321431087
RSME scores: [0.11690953 0.13573736 0.13798171 0.12709603 0.11467012]
Mean RSME: 0.12647895119428948

#fit LGBmodel
lgb_model.fit(X_train_scaled,y_train)

```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000501 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 2926
```

```
[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 65
```

```
[LightGBM] [Info] Start training from score 12.031035
```

```
LGBMRegressor(learning_rate=0.2, max_depth=6, n_estimators=50, num_leaves=6,
               reg_alpha=0)
```

```
# LGBmodel prediction
```

```
lgb_y_pred = lgb_model.predict(X_test_scaled)
```

```
# Calculate and print scores
```

```
lgb_r2_score = r2_score(y_test, lgb_y_pred)
```

```
lgb_mae_score = mean_absolute_error(y_test, lgb_y_pred)
```

```
lgb_rmse_score = np.sqrt(mean_squared_error(y_test, lgb_y_pred))
```

```
# Print scores
```

```
print("R2 score:", lgb_r2_score)
```

```
print("MAE scores:", lgb_mae_score)
```

```
print("RMSE score:", lgb_rmse_score)
```

```
R2 score: 0.9108859444420738
```

```
MAE scores: 0.08850765663252229
```

```
RMSE score: 0.1226512965590393
```

#####Voting Regressor

```
# Voting Model
```

```
voting_model = VotingRegressor(estimators=[('Gradient Boost', gbr_model), ('xgboost', xgb_model),
                                          ('lightgbm', lgb_model)])
```

```
# voting Hyperparameter
```

```
voting_parameters = {
    'weights': [[0.2,0.6,0.2], [0.6,0.2,0.2],[0.40,0.40,0.40], [0.20,0.40,0.50]]
}
```

```
voting_model_random = RandomizedSearchCV(voting_model,voting_parameters,n_iter=10, cv=5,scoring = 'neg_mean_absolute_error',n_jobs = -1)
```

```
voting_model_random.fit(X_train_scaled,y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:305: UserWarning: The total space of parameters 4 is smaller than n_iter=10. Running 4 iterations. For exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000465 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 2926
```

```
[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 65
```

```
[LightGBM] [Info] Start training from score 12.031035
```

```
RandomizedSearchCV(cv=5,
                   estimator=VotingRegressor(estimators=[('Gradient Boost',
                                                           GradientBoostingRegressor(alpha=0.5,
                                                           n_estimators=30)),
                                                           ('xgboost',
                                                           XGBRegressor(alpha=0.5,
```

```

base_score=None,
booster=None,
callbacks=None,
colsample_bylevel=None,
colsample_bynode=None,
colsample_bytree=0.7,
device=None,
early_stopping_rounds=None,

one,

e,

ne,

...)),

('lightgbm',
 LGBMRegressor(learning_rate=0.2,
                max_depth=6,
                n_estimators=50,
                num_leaves=6,
                reg_alpha=0))),

n_jobs=-1,
param_distributions={'weights': [[0.2, 0.6, 0.2],
                                [0.6, 0.2, 0.2],
                                [0.4, 0.4, 0.4],
                                [0.2, 0.4, 0.5]]},

scoring='neg_mean_absolute_error')

# figure out best hyperparameters
print(voting_model_random.best_estimator_)

VotingRegressor(estimators=[('Gradient Boost',
                             GradientBoostingRegressor(alpha=0.5,
                                                         n_estimators=30)),
                             ('xgboost',
                              XGBRegressor(alpha=0.5, base_score=None,
                                             booster=None, callbacks=None,
                                             colsample_bylevel=None,
                                             colsample_bynode=None,
                                             colsample_bytree=0.7, device=None,
                                             early_stopping_rounds=None,
                                             enable_categorical=False,
                                             eval_metric=None, feature_types=None,
                                             gamma=0.1,
                                             max_cat_threshold=None,
                                             max_cat_to_onehot=None,
                                             max_delta_step=None, max_depth=5,
                                             max_leaves=None,
                                             min_child_weight=1.7817, missing=nan,
                                             monotone_constraints=None,
                                             multi_strategy=None, n_estimators=220,
                                             n_jobs=None, num_parallel_tree=None, ...)),
                             ('lightgbm',
                              LGBMRegressor(learning_rate=0.2, max_depth=6,
                                             n_estimators=50, num_leaves=6,

```

```

        reg_alpha=0))],
        weights=[0.2, 0.4, 0.5])

# create model instance after hyperparameters tuning
voting_model = VotingRegressor(estimators=[('Gradient Boost', gbr_model), ('xgboost', xgb_model),
    ('lightgbm', lgb_model)], weights=[0.2, 0.4, 0.5])

# Validation of Voting Regressor
voting_r2_cv = cross_val_score(voting_model, X_train_scaled, y_train, scoring="r2", cv = 5)
voting_mae_cv = -cross_val_score(voting_model, X_train_scaled, y_train, scoring="neg_mean_absolute_error", cv = 5)
voting_rmse_cv = -cross_val_score(voting_model, X_train_scaled, y_train, scoring="neg_root_mean_squared_error", cv = 5)
print("R2 scores:", voting_r2_cv)
print("Mean R2:", np.mean(voting_r2_cv))
print("MAE scores:", voting_mae_cv)
print("Mean MAE:", np.mean(voting_mae_cv))
print("RSME scores:", voting_rmse_cv)
print("Mean RSME:", np.mean(voting_rmse_cv))

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000476 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2828
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.030302
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000446 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
[LightGBM] [Info] Start training from score 12.020671
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000471 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2822
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 64
[LightGBM] [Info] Start training from score 12.041320
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000441 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2824
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.034734
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000467 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2835
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.028148
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000550 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2828
[LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65

```

[LightGBM] [Info] Start training from score 12.030302
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000615 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2824
 [LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
 [LightGBM] [Info] Start training from score 12.020671
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000497 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2822
 [LightGBM] [Info] Number of data points in the train set: 934, number of used features: 64
 [LightGBM] [Info] Start training from score 12.041320
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000459 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2824
 [LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
 [LightGBM] [Info] Start training from score 12.034734
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000485 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2835
 [LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
 [LightGBM] [Info] Start training from score 12.028148
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000463 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2828
 [LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
 [LightGBM] [Info] Start training from score 12.030302
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000467 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2824
 [LightGBM] [Info] Number of data points in the train set: 934, number of used features: 65
 [LightGBM] [Info] Start training from score 12.020671
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000452 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2822
 [LightGBM] [Info] Number of data points in the train set: 934, number of used features: 64
 [LightGBM] [Info] Start training from score 12.041320
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000469 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 2824
 [LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
 [LightGBM] [Info] Start training from score 12.034734
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000511 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 2835
[LightGBM] [Info] Number of data points in the train set: 935, number of used features: 65
[LightGBM] [Info] Start training from score 12.028148
R2 scores: [0.90759188 0.87050146 0.89093514 0.88825343 0.90530335]
Mean R2: 0.892517051729099
MAE scores: [0.0799881 0.09111986 0.09358461 0.09280279 0.07613774]
Mean MAE: 0.08672662128709613
RSME scores: [0.11446415 0.13477808 0.13525746 0.12729988 0.10882885]
Mean RSME: 0.1241256842819197
```

Fit on data

```
voting_model.fit(X_train_scaled, y_train)
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000527 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 2926
```

```
[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 65
```

```
[LightGBM] [Info] Start training from score 12.031035
```

```
VotingRegressor(estimators=[('Gradient Boost',
                             GradientBoostingRegressor(alpha=0.5,
                                                         n_estimators=30)),
                             ('xgboost',
                              XGBRegressor(alpha=0.5, base_score=None,
                                             booster=None, callbacks=None,
                                             colsample_bylevel=None,
                                             colsample_bynode=None,
                                             colsample_bytree=0.7, device=None,
                                             early_stopping_rounds=None,
                                             enable_categorical=False,
                                             eval_metric=None, feature_types=None,
                                             gamma=0.0,
                                             max_cat_threshold=None,
                                             max_cat_to_onehot=None,
                                             max_delta_step=None, max_depth=5,
                                             max_leaves=None,
                                             min_child_weight=1.7817, missing=nan,
                                             monotone_constraints=None,
                                             multi_strategy=None, n_estimators=220,
                                             n_jobs=None, num_parallel_tree=None, ...)),
                             ('lightgbm',
                              LGBMRegressor(learning_rate=0.2, max_depth=6,
                                             n_estimators=50, num_leaves=6,
                                             reg_alpha=0))],
                weights=[0.2, 0.4, 0.5])
```

Voting prediction

```
voting_model_y_pred = voting_model.predict(X_test_scaled)
```

Calculate and print evaluation score

```
voting_r2_score = r2_score(y_test, voting_model_y_pred)
```

```
voting_mae_score = mean_absolute_error(y_test, voting_model_y_pred)
```

```
voting_rmse_score = np.sqrt(mean_squared_error(y_test, voting_model_y_pred))
```

Print scores

```
print("R2 score:", voting_r2_score)
```

```
print("MAE scores:", voting_mae_score)
```

```
print("RMSE score:", voting_rmse_score)
```

```
R2 score: 0.9122033750178448
MAE scores: 0.08632119223256815
RMSE score: 0.12174130416172735
```

###Summary


```
# summary result table
summary = pd.DataFrame({
    'Model Name': ['Lasso', 'Decision Tree', 'Support Vector Regressor', 'KNN', 'Random Forest',
    'Gradient Boosting', 'XGBoost', 'LightGBM', 'Voting Regressor'],
    'R2 Score': [lasso_r2_score, dtr_r2_score, svr_r2_score, knn_r2_score, rf_r2_score, gbr_r2_score,
    xgb_r2_score, lgb_r2_score, voting_r2_score],
    'MAE Score': [lasso_mae_score, dtr_mae_score, svr_mae_score, knn_mae_score, rf_mae_score, gbr_mae_score,
    xgb_mae_score, lgb_mae_score, voting_mae_score],
    'RMSE Score': [lasso_rmse_score, dtr_rmse_score, svr_rmse_score, knn_rmse_score, rf_rmse_score, gbr_rmse_score,
    xgb_rmse_score, lgb_rmse_score, voting_rmse_score]
})
df_summary = pd.DataFrame(summary)
df_summary
```

	Model Name	R2 Score	MAE Score	RMSE Score
0	Lasso	0.904482	0.089703	0.126982
1	Decision Tree	0.791153	0.144035	0.187764
2	Support Vector Regressor	0.821243	0.111036	0.173712
3	KNN	0.818839	0.128605	0.174876
4	Random Forest	0.889963	0.097287	0.136291
5	Gradient Boosting	0.877650	0.104677	0.143715
6	XGBoost	0.913757	0.085211	0.120659
7	LightGBM	0.910886	0.088508	0.122651
8	Voting Regressor	0.912203	0.086321	0.121741

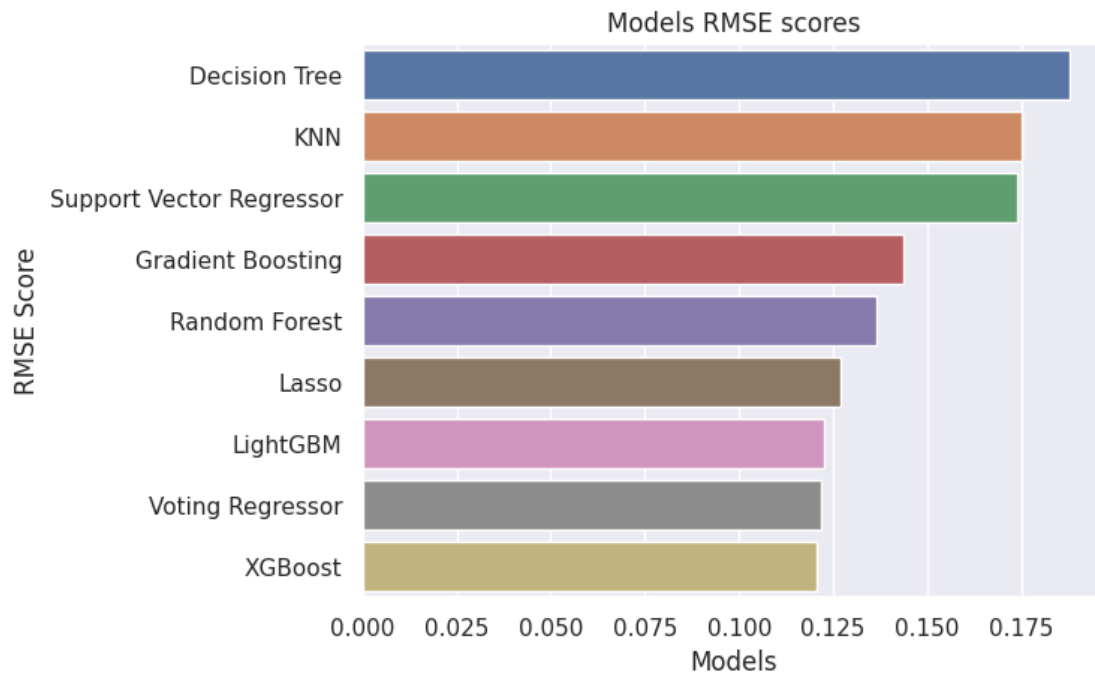
```
# Create Bar plot for models according to R2.
df_summary = df_summary.sort_values(by='R2 Score', ascending=False)
sns.barplot(y = df_summary['Model Name'], x = df_summary['R2 Score'], data= df_summary, orient
= 'h')
plt.title(f"Models R2 scores")
plt.xlabel('Models')
plt.ylabel("R2 Score")
plt.show()
```



```
# Create Bar plot for models according to MAE.
df_summary = df_summary.sort_values(by='MAE Score', ascending=False)
sns.barplot(y = df_summary['Model Name'],x = df_summary['MAE Score'] ,data= df_summary, orient
= 'h')
plt.title(f"Models MAE scores")
plt.xlabel('Models')
plt.ylabel("MAE Score")
plt.show()
```




```
# Create Bar plot for models according to RMSE.
df_summary = df_summary.sort_values(by='RMSE Score', ascending=False)
sns.barplot(y = df_summary['Model Name'], x = df_summary['RMSE Score'], data= df_summary, orient = 'h')
plt.title(f"Models RMSE scores")
plt.xlabel('Models')
plt.ylabel("RMSE Score")
plt.show()
```



```
#XGBoost Feature importance
for score, name in zip(xgb_model.feature_importances_, X_train_scaled):
    print(round(score,2),name)
```

```
0.0 MSSubClass
0.01 MSZoning
0.0 LotFrontage
0.0 LotArea
0.0 Street
0.0 Alley
0.0 LotShape
0.01 LandContour
0.0 Utilities
0.0 LotConfig
0.0 LandSlope
0.0 Neighborhood
0.0 Condition1
0.0 Condition2
0.0 BldgType
0.0 HouseStyle
0.33 OverallQual
0.01 OverallCond
0.03 YearBuilt
0.01 YearRemodAdd
0.0 RoofStyle
0.0 RoofMatl
0.0 Exterior1st
```

0.0 Exterior2nd
0.0 MasVnrType
0.0 MasVnrArea
0.14 ExterQual
0.0 ExterCond
0.0 Foundation
0.01 BsmtQual
0.0 BsmtCond
0.0 BsmtExposure
0.0 BsmtFinType1
0.01 BsmtFinSF1
0.0 BsmtFinType2
0.0 BsmtFinSF2
0.0 BsmtUnfSF
0.02 TotalBsmtSF
0.0 Heating
0.0 HeatingQC
0.03 CentralAir
0.0 Electrical
0.01 1stFlrSF
0.01 2ndFlrSF
0.0 LowQualFinSF
0.05 GrLivArea
0.0 BsmtFullBath
0.0 BsmtHalfBath
0.02 FullBath
0.01 HalfBath
0.0 BedroomAbvGr
0.0 KitchenAbvGr
0.05 KitchenQual
0.0 TotRmsAbvGrd
0.0 Functional
0.03 Fireplaces
0.0 FireplaceQu
0.02 GarageType
0.0 GarageYrBlt
0.01 GarageFinish
0.11 GarageCars
0.01 GarageArea
0.01 GarageQual
0.02 GarageCond
0.0 PavedDrive
0.0 WoodDeckSF
0.0 OpenPorchSF
0.0 EnclosedPorch
0.0 3SsnPorch
0.0 ScreenPorch
0.0 PoolArea
0.0 PoolQC
0.0 Fence
0.0 MiscFeature
0.0 MiscVal
0.0 MoSold
0.0 YrSold
0.0 SaleType
0.0 SaleCondition